DETECTION OF FALLEN ELECTRIC CABLESON ROAD AND STREETS IN INNER CITIES FOR AUTONOMOUS VEHICLES AND BIKES

UBRAN FALLEN CABLE DETECTION DATASET (UFCDD): A VISUAL BENCHMARK FOR SAMRT MOBILITY IN PAKISTAN

Humza Nazeer

Humza Nazeer
(MUST/FA22-RCS-009/AJK)

Session 2022-2024

Department of CS & IT
Faculty of Natural & Applied Sciences

Mirpur University of Science and Technology (MUST), Mirpur
AJK Pakistan

DETECTION OF FALLEN ELECTRIC CABLESON ROAD AND STREETS IN INNER CITIES FOR AUTONOMOUS VEHICLES AND BIKES

UBRAN FALLEN CABLE DETECTION DATASET (UFCDD): A VISUAL BENCHMARK FOR SAMRT MOBILITY IN PAKISTAN

By

Humza Nazeer

(MUST/FA22-RCS-009/AJK)

A thesis is submitted in partial fulfillment of
the requirement for the degree of

MASTER OF PHILOSOPHY

In

COMPUTER SCIENCE

Session 2022-2024

Department of CS& IT
Faculty of Natural & Applied Sciences

Mirpur University of Science and Technology (MUST), Mirpur (AJK) Pakistan

CERTIFICATION

I solemnly declare that this research is my own work, and it is an original one, and no part of this thesis falls under plagiarism. If found otherwise, at any stage, I will be responsible for the consequences.

Student's Name: **Humza Nazeer** Signature: _____

Registration No.: **MUST/FA22-RCS- 009/AJK** Date: _____

Certified that the contents and form of the thesis entitled **"DETECTION OF FALLEN ELECTRIC CABLESON ROAD AND STREETS IN INNER CITIES FOR AUTONOMOUS VEHICLES AND BIKES"** submitted by "**Humza Nazeer**" have been found satisfactory for the requirement of the degree.

Supervisor: Samreen Ayaz

Co-Supervisor: _____

External Examiner: _____

Date of Viva Voce: _____

Chairperson: _____

Dean of Faculty: _____

Director, Advanced Studies & Research Board: _____

DEDICATION

This work is dedicated to my family, especially my parents, wife, and my beloved kids for continuous support and prayers.

CONTENT

ABSTRACT* 10
Chapter No. 01 12
INTRODUCTION 12
1.1 Background 12
1.2 Problem Statement 14
1.3 Research Questions 15
1.4 Research Contributions 16
1.4.1 Creation and Dissemination of a Novel Dataset 16
1.4.2 Implementation of a Specialized High-Performance Detection Framework ... 17
1.4.3 Quantitative Validation of Robust Thin-Structure Detection 17
1.5. Research Objectives 18
1.6. Research Gap 19
1.6.1 Inherent Geometric and Photometric Challenges in Detection 19
1.6.2 Deficiencies in Standard Autonomous Vehicle Datasets 20
1.6.3 The Domain Transfer and Infrastructural Gap 21
Chapter No. 02 23
RELATED WORK 23
2.1. Foundations of Autonomous Vehicle Perception and System Architecture 23
2.1.1. Perception Requirements and Sensor Modalities 23
2.1.2 Sensor Limitations in Detecting Thin Structures 24
2.2 The Critical Challenge: High-Stakes, Low-Frequency (HSLF) Obstacles 24
2.2.1 Defining Unconventional Road Hazards and Infrastructure Failures 24
2.2.2 The Geometric Ambiguity Problem 25
2.3 Evolution of Hazard Detection Methodologies 27
2.3.1 Classical Computer Vision Approaches 27
2.3.2 The Paradigm Shift to Deep Learning 28
2.4 Deep Learning Architectures for Real-Time Object Detection in AVs 28
2.4.1 Taxonomy of Modern Detectors and the YOLO Framework 28
2.4.2 Hybrid Models and Lightweight Backbones for Edge Computing 29

2.5 Specialized Techniques for Thin and Curvilinear Structure Perception 32
2.5.1 Representation and Boundary Accuracy 32
2.5.2 Contextual and Depth Requirements for Safety 33
2.6 Addressing Data Scarcity and Training Difficulties 33
2.6.1 The Problem of Extreme Class Imbalance (EVI) 34
2.6.2 Specialized Loss Functions for Robust Learning 34
2.6.3 Synthetic Data Generation (SDG) 34
2.7 Analysis of Data Deficiencies and Domain Shift 35
2.7.1 Critical Review of Standard AV Benchmarks 35
2.7.2 Existing Datasets for Debris and Unstructured Ground 35
2.8 Conclusion and Synthesis: Paving the Path for Safer Smart Mobility 38
Chapter No. 03 40
MATERIALS AND METHODS 40
3.1 Data Acquisition and the Urban Fallen Cable Detection Dataset (UFCDD) 40
3.1.1 Image Acquisition Platform and Equipment 41
3.1.2 Environmental and Temporal Diversity 42
3.2 Dataset Preparation and Ground Truth Annotation 43
3.2.1 Preprocessing and Key Frame Selection 43
3.2.2 Manual Annotation Process 43
3.2.3 Dataset Partitioning 44
3.3 Specialized Model Architecture and Optimization 45
3.3.1 YOLOv5 as the Foundational Architecture 45
3.3.2 Custom Architectural Configuration 46
3.4 Training Strategy and Experimental Setup 47
3.4.1 Computational Environment 47
3.4.2 Training Hyperparameters 48
3.4.3 Loss Function Configuration 49
Chapter No. 04 51
RESULTS AND DISCUSSION 51
4.1 Quantitative Performance Results on the UFCDD 51

4.2 Discussion of Core Findings: Classification vs. Localization 54
4.2.1 Success in Hazard Classification (mAP@50) 54
4.2.2 Localization Deficit and the Geometric Challenge (mAP@50–95) 58
4.3 Comparative Robustness and Architectural Validation 60
4.4 System Efficiency and Real-Time Feasibility (RQ2 & RQ4) 63
4.5 Evaluation of Research Questions 64
Chapter No. 05 65
CONCLUSION 65
5.1 Future Work 66
5.1.1 Resolving Geometric and Training Limitations 66
5.1.2 Enhancing Robustness and Scalability 67
References: 68

LIST OF TABLES

Table No.	Page No.
3.1 Real Road Traffic Parameters	14
3.2 PAGE of Proposed Agent	17
4.2 Experiment Type_ A Set 1: Parameters and Their Values	24
4.3 Experiment Type_ A Set 2: Parameters and Their Values	24
4.4 Experiment Type_ A Set 3: Parameters and Their Values	25
4.5 Experiment Type_ A Set 4: Parameters and Their Values	25
4.6 Experiment Type_ A Set 5: Parameters and Their Values	26
4.7 Set 1 Type A Experiment Results	27
4.8 Set 2 Type A Experiment Results	27
4.9 Set 3 Type A Experiment Results	29
4.10 Set 4 Type A Experiment Results	29
4.11 Set 5 Type A Experiment Results	31

COPY RIGHTS

(ADD THIS PAGE, IF YOU HAVE ORGINAL RESEARCH DATA WHICH HAVE NOT BEEN YET PUBLISHED)

THE DATA AND RESULTS REPORTED IN THE TEHSIS HAS NOT BEEN PUBLSIHED OR SHARED YET AND PROTECTED BY COPY RIGHTS OWNED BY UNIVERSITY AND PROHIBITED TO USE FOR ANY PRUPOSE WITHOUT PERMISSION.

Acknowledgments

I am thankful to Allah who blessed me with the courage and potential to perform this work. With His kindness, I have been able to accomplish this task. I am also thankful to my parents, my wife, and my kids for their continuous prayers and support. I am thankful to my supervisor Dr Samreen Ayaz, whose supervision and direction made it possible for me to complete this work. Throughout the research work, his motivation and guidance helped me.

Humza Nazeer

ABSTRACT*

Abstract

Autonomous vehicles (AVs) rely critically on robust perception systems to ensure safe operation in complex urban environments. While contemporary perception frameworks demonstrate strong performance for common road entities, they remain vulnerable to High-Stakes, Low-Frequency (HSLF) hazards, particularly thin, irregular objects such as fallen electric and communication cables. These hazards pose severe safety risks, including vehicle damage, electrical failure, and potential loss of control, and are especially prevalent in unstructured urban settings of developing regions where infrastructure maintenance is inconsistent. Existing autonomous driving datasets and detection models largely overlook such hazards, resulting in poor generalization and high false-negative rates. This research addresses this critical safety gap by proposing a dedicated deep learning–based detection framework for fallen cable hazards. First, a novel dataset, the **Urban Fallen Cable Detection Dataset (UFCDD)**, is developed, comprising over 2,500 carefully annotated images collected from real-world urban environments in Pakistan under diverse lighting and weather conditions. Second, a specialized Convolutional Neural Network (CNN) based object detection model is implemented using the YOLOv5 architecture, optimized for real-time detection of thin, curvilinear structures in cluttered scenes. The framework classifies cable conditions as safe or unsafe and emphasizes precise localization to support downstream navigation and path planning. Extensive experimental evaluation demonstrates that the proposed approach achieves a mean Average Precision of 96% at mAP@50 and 76% at the more stringent mAP@50–95, confirming its effectiveness in both detection and fine-grained localization of low-aspect-ratio hazards. These results represent a significant improvement over generalized perception pipelines and validate the feasibility of reliable thin-structure detection in safety-critical autonomous driving contexts. Overall, this work contributes a novel dataset, a validated real-time detection framework, and empirical evidence supporting enhanced AV safety and scalability, particularly for

deployment in developing countries with unstructured and infrastructure-challenged urban environments.

Chapter No. 01

INTRODUCTION

1.1 Background

1.1.1 The Advancement of Autonomous Mobility and Perception Systems

The past several years has witnessed a decisive focus on enhancing technology to expedite tasks and improve human life, culminating in the development of the autonomous vehicle (AV), which is now considered a principal component of modern transportation systems. Autonomous vehicles, classified across six levels of automation (SAE Levels 0–5), promise significant improvements in mobility, efficiency, and public safety. The imperative for this technology is underpinned by compelling safety statistics: human error contributes to the vast majority of traffic incidents, resulting in approximately 1.35 million annual deaths globally according to World Health Organization reports[1]. By minimizing human factors such as fatigue, distraction, or impaired driving, AVs offer the potential to drastically reduce these catastrophic failures. Furthermore, the deployment of AVs enhances inclusivity by providing mobility to populations unable to drive and contributes to environmental sustainability through optimized traffic flow and reduced emissions [1].

The operational integrity of an AV hinges upon its perception system, which is tasked with the real-time identification, classification, and localization of objects in the surrounding environment [2]. This complex functionality relies heavily on sensor fusion, integrating data from multiple modalities. Cameras supply rich visual and semantic information regarding color, texture, and lane markings. LiDAR provides high-fidelity three-dimensional (3D) point clouds critical for depth and spatial understanding, aiding in localization and mapping. Radar systems offer robust distance and speed measurements, proving invaluable during adverse weather where optical sensors struggle. These inputs are processed through sophisticated deep learning frameworks, typically Convolutional Neural Networks (CNNs) and CNN-Transformer hybrids, which

[1] https://www.who.int/data/gho/data/themes/topics/topic-details/GHO/road-traffic-mortality

automatically extract hierarchical features from raw data to identify objects like vehicles, pedestrians, and traffic signs [3].

Despite the impressive strides made by state-of-the-art perception systems, particularly in detecting common, structured road elements, performance degradation occurs when these systems encounter geometrically complex, irregular, or unconventional hazards [4]. A notable constraint for current sensor fusion architectures lies in their design efficiency; while sensors provide redundancy for bulky objects, for example, radar compensating for camera failures in fog, they exhibit limitations against fine structures. LiDAR often yields sparse point cloud returns for thin objects, and radar cross-sections are too low for reliable detection. This forces the perception framework to rely almost entirely on high-resolution monocular vision, demanding that the accompanying detection algorithm achieve exceptional precision and robustness to background clutter to compensate for the lack of reliable complementary depth data from other sensors. The successful deployment of AVs worldwide thus requires the perception of systems to be robust against unpredictable environments and infrastructure failures, leading to dedicated research into these overlooked hazard categories.

1.1.2 **The High-Stakes, Low-Frequency Hazard of Fallen Cables**

Fallen electric and communication cables on roadways represent a specific, critical category of High-Stakes, Low-Frequency (HSLF) hazards. The implications of failing to detect these objects are severe; running over them can cause significant mechanical damage to the vehicle's undercarriage, lead to total system failure due to electrical shorts, and potentially result in loss of control, endangering occupants and the public [5]. The fatality rates associated with power line incidents in other transportation domains, such as the 30% fatality rate reported by the Federal Aviation Administration for aircraft-power line collisions between 1997 and 2010, highlight the extreme danger associated with infrastructure failures involving cables.

This challenge is exacerbated in unstructured urban environments, particularly in developing regions, where infrastructure maintenance is often irregular and density is high. The research material specifically references Pakistan, mentioning that such environments present unique operational challenges not typically covered by Western-focused global benchmarks [6]. Moreover, environmental factors, including the increased frequency of severe weather events like

dust storms, a consequence of climate change, contribute directly to electrical cables breaking and falling onto the roads [7]. To achieve global safety certification and scalable deployment, AVs must demonstrate reliable performance against these regionally specific infrastructural and climatic failures.

1.2 Problem Statement

The absence of reliable, real-time detection mechanisms for fallen electric cables in high-clutter urban environments constitutes a fundamental and critical safety risk that threatens the trustworthiness and global scalability of autonomous vehicle technology.

Fallen and low-hanging electric cables pose a critical threat, serving as dangerous road hazards that can obstruct traffic, damage sophisticated vehicle components, and pose electrocution risks to both humans and the vehicle's electrical systems [5]. Existing perception systems fail due to the inherent complexity of these objects; they are not readily visible to generalized vision apparatus, particularly within the complex and dynamic visual background of an inner city [7]. Geometrically, these cables are defined by a critically low aspect ratio, often spanning only a few pixels in a camera image. This results in visual ambiguity, causing them to be confused with background noise, shadows, or benign linear markings, leading to high false negative rates where the hazard is missed entirely.

This technological shortcoming is compounded by a critical deficiency in training resources. Standard autonomous vehicle datasets, which overwhelmingly prioritize common traffic features such as pedestrians, vehicles, and road signs, completely overlook thin, irregular hazards like fallen electric cables [8]. Consequently, specialized deep learning models capable of identifying the fine-grained, curvilinear features of these objects cannot be adequately trained and validated. This lack of specific data and tailored detection models for thin, irregular hazards represents a significant safety gap.

Therefore, the objective is to develop a robust, deep learning-based detection framework, supported by a novel, specialized dataset, the Urban Fallen Cable Detection Dataset (UFCDD),

to effectively identify these risks and ensure safe autonomous vehicle navigation in complex, unstructured urban areas.

1.3 Research Questions

The following research questions are designed to address the challenges outlined in the Problem Statement by guiding the systematic development of the requisite data and detection framework.

RQ1: Data Generation and Fidelity: How can a dedicated, high-fidelity visual dataset (UFCDD) be systematically collected and annotated to effectively capture the diverse geometric and photometric properties of fallen and low-hanging electric cables across varied environmental (clear, overcast) and lighting conditions (dawn, dusk, day) in unstructured urban settings?

This question establishes the foundational necessity of creating a dataset that accurately reflects the real-world complexity and variability of cable hazards in the target domain, which previous benchmarks failed to achieve [1].

RQ2: Model Architecture and Optimization: Can a Convolutional Neural Network (CNN)-based object detection framework, specifically tailored using the YOLOv5 architecture, be robustly trained on the specialized Urban Fallen Cable Detection Dataset (UFCDD) to achieve real-time, high-accuracy detection of thin, curvilinear cable hazards?

This question explores the viability of CNNs, known for their feature extraction capabilities, in resolving the technical difficulty posed by thin structures. The selection of YOLOv5, a state-of-the-art single-stage detector, is implicitly critical, as AV systems require minimal latency for safety-critical maneuvers [2]. The architectural premise is that YOLOv5 offers the necessary balance of speed and precision, essential for translating detection into actionable, real-time safety measures.

RQ3: Performance Validation and Localization Precision: What level of real-time performance (quantified by metrics such as mAP@50 and the more stringent mAP@50-95) can the specialized YOLOv5 framework achieve in accurately localizing and classifying fallen cable

hazards, thereby demonstrating a tangible advancement in robustness against the geometric challenges of thin structures?

This focuses on rigorous empirical validation. Since thin objects are hard to localize precisely, assessing the model against stringent metrics like mean Average Precision at Intersection over Union thresholds from 50% to 95% (mAP@50-95) is essential to confirm not merely detection, but accurate boundary definition, which is necessary for effective path planning around the hazard . The localization accuracy confirms the model's ability to overcome the geometric ambiguity inherent in thin, low-pixel-count objects.

RQ4: Impact on Safety and Scalability: How does the proposed detection framework, enabled by the UFCDD, enhance autonomous vehicle safety and reliability in complex urban environments, and what are the implications for scalable smart mobility solutions in developing countries where these hazards are prevalent?

This question addresses the practical application and broader relevance of the research, linking the technical solution to societal benefits, especially in environments globally representative of varying infrastructural quality [6].

1.4 Research Contributions

The research delivers several specialized contributions that advance state-of-the-art in autonomous vehicle perception and safety, particularly concerning unconventional hazards in unstructured environments.

1.4.1 Creation and Dissemination of a Novel Dataset

The primary contribution is the development and collection of the Urban Fallen Cable Detection Dataset (UFCDD). This dataset comprises a novel collection of over 2,500 meticulously annotated images specifically capturing fallen and low-hanging electric and communication cables in urban settings, primarily sourced from Mirpur, AJK, Pakistan. This collection directly addresses a critical oversight in existing datasets, which traditionally focus on vehicles, pedestrians, and road signs, while ignoring these cable-induced hazards [6]. The UFCDD is structured to include images captured under diverse conditions, early morning, afternoon, and evening, to ensure the resulting models are robust to varying illumination and weather, mitigating common performance

constraints in AV systems. By focusing on a region outside the typical Western-centric AV testing domain, the UFCDD provides necessary geographical and infrastructural diversity, essential for developing AV systems capable of operating safely on a global scale.

1.4.2 Implementation of a Specialized High-Performance Detection Framework

A specialized deep learning object detection framework, leveraging the efficiency and accuracy of the YOLOv5 architecture, is a core contribution. This framework is premised on the capabilities of Convolutional Neural Networks (CNNs) to automatically learn and extract a hierarchy of feature representations from raw image data, a capability vital for identifying the space-based fine-grained features (edges, curves, and shapes) required for recognizing thin, elongated objects in complex urban clutter . YOLOv5, selected as a state-of-the-art single-stage detector, balances detection accuracy with computational speed, crucial for the low-latency inference requirements of safety-critical autonomous driving [1, 7, 8, 9, 10]. The model is trained to differentiate between "Safe" (intact) and "Unsafe" (fallen or hazardous) cable conditions, translating complex visual information into direct, navigational decision-making inputs.

1.4.3 Quantitative Validation of Robust Thin-Structure Detection

The research successfully validated the performance of the specialized CNN-based YOLOv5 model, providing quantifiable evidence of its efficacy in detecting thin, irregular hazards. The model achieved a 96% accuracy at mAP@50 and a 76% accuracy at mAP@50-95. The achievement of a high mAP@50 score indicates strong classification capability, reliably distinguishing hazardous cables from the background. Furthermore, the robust 76% accuracy at the mAP@50-95 threshold is particularly significant. Since this metric requires a highly precise boundary definition (high Intersection over Union scores), the result confirms that the framework successfully overcame the fundamental geometric challenge posed by low-pixel-count objects. This successful validation demonstrates the capability of the tailored model to identify thin, irregular cable structures even in difficult urban conditions, contributing to a reliable detection system that enhances urban transportation safety [13].

Table 1.1: Key Contributions and Performance Benchmarks

Contribution Element	**Description/Metric**	**Significance**
Dataset Name	Urban Fallen Cable Detection Dataset (UFCDD)	Addresses overlooked High-Stakes, Low-Frequency (HSLF) hazards [14].
Dataset Size/Scope	2,500+ images from Mirpur, AJK, Pakistan	Provides critical geographical and infrastructural diversity for global AV deployment [15].
Detection Framework	CNN-based YOLOv5 Model	Ensures real-time, low-latency performance required for safety-critical AV tasks [16].
Mean Average Precision (mAP@50)	96%	High accuracy for primary detection threshold, indicating strong capability in distinguishing hazardous objects [17].
Mean Average Precision (mAP@50-95)	76%	Robust performance across stricter localization tolerances, confirming precision on thin, complex structures [18].

1.5. Research Objectives

The research is guided by a set of concrete objectives designed to systematically address the identified safety gap through data creation and model implementation.

1. **Hazard Image Collection:** To collect a comprehensive set of images capturing fallen power and communication cables from various urban locations in Pakistan, ensuring the data reflects real-world infrastructural diversity. The scope of this collection is defined to include multiple areas in Pakistan (Mirpur) to account for geographical and infrastructural variations.
2. **Environmental Diversity Acquisition:** To gather data across a full spectrum of natural lighting and weather conditions, encompassing clear daylight, overcast/cloudy weather, and specific time intervals such as early morning (dawn), late morning, afternoon, and evening (dusk). This ensures that the model is trained against the visual variations common in dynamic environments.
3. **Manual Annotation and Ground Truth Generation:** To manually tag the collected imagery with high precision, specifically marking the spatial boundaries of fallen cables to create reliable ground truth. This process is critical for training the supervised deep learning model. Public
4. **Dataset Construction:** To structure, validate, and build the Urban Fallen Cable Detection Dataset (UFCDD) as a publicly accessible resource to help train and test hazard detection systems for autonomous vehicles and broader robotics applications, including drone navigation and surveillance systems.
5. **Model Development and Validation:** To train a CNN-based object detection model (YOLOv5) on the UFCDD and quantitatively evaluate its real-time performance using widely recognized metrics such as mAP@50, mAP@50-95, precision, recall, and F1-score, thereby confirming its suitability for low-latency, safety-critical AV operation [11, 12, 13].

1.6. Research Gap

The justification for this specialized research lies in addressing several critical gaps that collectively limit the safety and reliability of autonomous vehicles, particularly in globally diverse and unstructured urban environments.

1.6.1 Inherent Geometric and Photometric Challenges in Detection

The core technical gap stems from the unique physics and geometry of fallen cables, which fundamentally challenge conventional computer vision systems. Cables are characterized by a critically low aspect ratio, often manifesting as sparse lines of only 1 to 3 pixels in camera images,

making them functionally invisible or ambiguous to generalized object detectors. Traditional bounding box frameworks, optimized for larger targets like vehicles, lack the necessary fidelity to accurately enclose or classify these sub-scale, thin objects. Furthermore, the visual scene is characterized by high photometric ambiguity. Fallen cables frequently exhibit low contrast against complex road surfaces (concrete, asphalt, dirt) and blend seamlessly with background noise, shadows, and innocuous linear features. This difficulty requires the specialized extraction of fine-grained features, such as subtle edges and curves, a task that standard perception pipelines are not optimized to perform [18].

1.6.2 Deficiencies in Standard Autonomous Vehicle Datasets

A primary barrier to developing a robust detection solution is the pervasive data deficiency in established benchmarks. Standard AV datasets, including major public resources like KITTI, NuScenes, and BDD100K, are insufficient for this task. These benchmarks were overwhelmingly collected in highly structured, developed urban or highway settings, inherently minimizing the occurrence of unpredictable infrastructure failures like widespread debris or utility disruptions. The research gap is defined by two key failures in existing data:

1. **Annotation Scarcity:** These popular datasets lack the specific taxonomy and precise annotation necessary for thin, unconventional hazards. This absence prevents the development of machine learning models capable of generalizing dynamic, high-clutter environments.
2. **Misaligned Debris Focus:** While some specialized datasets exist for generalized road debris (e.g., HazardNet), they often focus on larger, non-curvilinear obstacles, such as tires or construction materials, frequently employing height thresholds that explicitly exclude the critically low-lying geometry of fallen electrical lines.

The lack of specialized training data forces models to rely on features learned from irrelevant or benign contexts, resulting in poor detection performance for HSLF cable hazards. The creation of the UFCDD directly addresses this data gap by curating imagery specific to high-risk, unstructured urban road conditions.

1.6.3 The Domain Transfer and Infrastructural Gap

A critical operational gap exists between successful prior research and the current requirement for ground-based autonomous vehicles. Earlier academic efforts related to line detection focused heavily on the aerial inspection of overhead power lines and pylons using Unmanned Aerial Vehicles (UAVs) [10]. These methodologies employed techniques like Radon transform, Kalman filtering, and feature grouping, which performed adequately in controlled environments with predictable backgrounds [11]. However, the transfer of these aerial techniques to ground AVs faces severe limitations. Aerial systems benefit from high contrast (wire against a clear sky) and structured, straight-line geometries. Ground AVs must instead contend with high background clutter, complex road textures, visual interference from rainfall, and the highly complex, non-rigid, and curved paths of fallen cables. These environmental and topological distinctions mandate fundamental adaptation of detection strategies. Furthermore, the focus on unstructured urban environments in regions like Pakistan highlights a fundamental geopolitical and infrastructural gap in AV safety development. The prevalence of fallen cables in such areas is not a random anomaly, but a predictable risk linked to factors such as infrastructure density, irregular maintenance, and environmental stress [6]. The requirement for a region-specific dataset like the UFCDD demonstrates that global AV maturity (SAE Level 5) cannot be achieved solely through centralized, generalized models, but necessitates modular perception stacks capable of reliable domain adaptation to diverse, real-world infrastructure quality and regional hazards. This research, by providing the necessary foundation for recognizing and reacting to these specialized regional failures, plays a vital role in enabling scalable smart mobility solutions in developing countries.

Chapter No. 02: Relevant Research in This Field

Chapter No. 3: Describe the materials and experimental configurations in detail.

Chapter 4: Results of suggested setups and study.

Chapter No. 5: How to Expand This Research and Potential Future Directions for Research Extension.

Chapter No. 02

RELATED WORK

2.1. Foundations of Autonomous Vehicle Perception and System Architecture

The transition toward fully autonomous mobility (SAE Levels 4 and 5) necessitates the development of highly reliable perception systems capable of achieving comprehensive environmental understanding under all operational conditions.[1] Autonomous vehicles (AVs) operate based on an interconnected stack of modules, including perception, localization, planning, and control, where the perception system provides the essential, safety-critical data input for all subsequent decision-making processes.[2]

2.1.1. Perception Requirements and Sensor Modalities

Robust perception systems in modern AVs rely on a multi-modal sensor suite to compensate for the inherent limitations of any single sensor type.[3]

The Visual Modality (Cameras) is indispensable, providing rich semantic information, color, texture, and lane markings, that is crucial for complex scene comprehension.[1] Camera data feeds into sophisticated deep learning architectures, such as You Only Look Once (YOLO)[42] and Faster R-CNN,[13] to perform object detection in dynamic urban environments.[1] Although cost-effective and simple to implement, vision systems are highly susceptible to performance degradation during challenging situations, including low light, heavy glare, fog, or rain, which compromise image quality and system reliability.[7]

Ranging Modalities, primarily Light Detection and Ranging (LiDAR) and Radar, supplement visual data. LiDAR generates high-fidelity three-dimensional (3D) point clouds, providing accurate distance, geometric, and spatial information essential for precise localization and obstacle avoidance.[1] Radar, conversely, excels in determining distance and speed, offering superior robustness in adverse weather conditions (rain, fog, dust) where optical sensors often fail [4].

2.1.2 Sensor Limitations in Detecting Thin Structures

A critical challenge arises when these standard sensor modalities encounter geometrically complex, thin structures, such as fallen electric cables. The inherent properties of range-based sensors render them unreliable for fine wire detection.

For LiDAR, the critically low aspect ratio of a cable (its thin profile relative to its length) results in a sparse point cloud return.[5] This lack of geometric information makes it difficult for LiDAR processing algorithms to reliably identify the obstacle.[5] Furthermore, radar's primary limitation is its inherently low resolution, which provides insufficient detail to detect small, static, or fine obstacles, such as fallen wires, making it inadequate for this specific hazard recognition task.[5]

The inability of ranging sensors to reliably identify these hazards places a disproportionate burden on the vision-based system.[5] Consequently, successful detection must rely overwhelmingly on high-resolution visual processing systems and specialized deep learning models to extract the fine semantic features (edges, shading, texture) of the cable structure where geometric sensors fail.[5] This necessitated research focus drives the architectural constraints discussed in later sections.

2.2 The Critical Challenge: High-Stakes, Low-Frequency (HSLF) Obstacles

While current state-of-the-art AV perception systems are highly proficient in identifying common, structured traffic elements, vehicles, pedestrians, and traffic signs, they exhibit significant limitations when faced with rare, geometrically complex, or unconventional road hazards, collectively categorized as road debris or infrastructure failures.[1]

2.2.1 Defining Unconventional Road Hazards and Infrastructure Failures

Fallen electric and communication cables on urban roadways represent a specific, crucial category of High-Stakes, Low-Frequency (HSLF) hazards.[1] The potential risks associated with running over such an obstacle are catastrophic, including severe mechanical damage to the vehicle's undercarriage, system failure due to electrical shorts, and loss of control, which endangers both vehicle occupants and the public.[1] Agencies such as the Federal Aviation Administration (FAA) and the American Automobile Association (AAA) have underscored the severity of collisions and

crashes linked to utility lines and road debris, validating the necessity of robust detection systems.[17]

This safety problem is particularly acute in unstructured urban environments, such as those found in developing regions like Mirpur, Pakistan.[25] In these settings, high infrastructure density, irregular maintenance, and increased vulnerability to environmental factors, specifically the heightened frequency of severe weather events like dust storms leading to cable breakage, exacerbate the risk of electrical cables falling onto the roadway.[25] For autonomous mobility to achieve global, scalable deployment, systems must possess the capability to recognize and react reliably to these region-specific infrastructural failures.[25]

2.2.2 The Geometric Ambiguity Problem

Fallen electric cables present unique geometric and photometric challenges that distinguish them from traditional, large obstacles. Geometrically, they are defined by an extremely low aspect ratio, meaning they are very thin relative to their length, often spanning only a marginal number of pixels (sometimes 1 to 3 pixels) in high-resolution camera images, especially when viewed from a distance.[43]

This low pixel count creates profound visual ambiguity and clutter. Thin objects are easily confused with irrelevant background noise, dense road textures, cast shadows, or benign linear features like road cracks or painted markings.[45] The visual appearance (contrast, texture, color) of the cable relative to the road surface changes drastically depending on the time of day or the weather, resulting in low-contrast situations that confound generalized detectors.[45]

The implication of this geometric complexity is that models trained on structured datasets, which implicitly learn to disregard high levels of visual clutter, suffer an acute performance degradation, a phenomenon known as domain shift, when deployed in a high-clutter, unstructured environment.[45] The model, having been trained on clean data, exhibits catastrophic generalization failure where the rare, yet critical, HSLF hazard is ignored entirely [45].

Table 2.1 Thin-Object, Boundary-Aware, Powerline & Infrastructure Detection Research

Ref No.	Area / Focus	Key Findings	Relevance to Fallen Cable Detection
[9]	Aerial Power Line Extraction	Edge + region-based extraction	**Foundational thin-line detection** basis
[18]	Automated Power Line Inspection (Review)	**DL approaches** for line inspection	Techniques **transferable** to fallen cables
[36]	Power-line Segmentation	Semantic segmentation of overhead lines	**Directly adaptable** for fallen cable segmentation
[44]	Synthetic Wire Detection	**Synthetic → real** training improves robustness	Useful for **rare cable-fall training data**
[46]	Wire Detection + Depth	Cable and **depth estimation**	Useful for **locating** fallen wires on road
[51]	Small-object Sensitive Segmentation	**Attention** boosts thin-object sensitivity	**Critical** for long thin cable structures

[50]	Boundary-aware Loss Functions	Better **segmentation edges**	Helps separate cable pixels cleanly
[58]	**Thin-structure Obstacle Detection**	**Fast detection** of thin rods/wires	**Closest match** to fallen cable detection problem
[59]	Powerline & Pylon Detection	**Robust cable detection** from aerial views	**Cable shape priors** applicable to street-level detection

2.3 Evolution of Hazard Detection Methodologies

The necessity of reliably detecting thin, complex structures on the ground has driven a technological evolution, marking a fundamental shift from classical computer vision to contemporary deep learning techniques.

2.3.1 Classical Computer Vision Approaches

Early research in utility infrastructure focused predominantly on the inspection and detection of overhead power lines and pylons using imagery captured by unmanned aerial vehicles (UAVs).[1] These classical methodologies relied heavily on mathematically intensive, hand-crafted feature extraction. Key techniques included the Radon Transform and the Hough Transform, which were employed to detect and reconstruct straight line segments based on predictable geometric properties.[48] Techniques such as Kalman Filtering were subsequently used to track these reconstructed line structures across sequential frames, aiding in inspection.[1]

However, these traditional methods proved highly fragile and non-robust when applied to the ground-based domain of autonomous vehicles. They were highly sensitive to background interference, shadows, and the texture variability of urban roads.[1] Most critically, they lacked the generalization ability necessary to model the complex, non-rigid, and curved paths

characteristic of fallen cables on the ground, a topology that differs markedly from the structured, taut straightness of overhead power lines.[1] This failure to generalize in complex ground clutter mandated the adoption of a more robust methodological paradigm.[1]

2.3.2 The Paradigm Shift to Deep Learning

The inherent limitations of classical approaches catalyzed a decisive shift toward deep learning (DL) techniques.[54] Convolutional Neural Networks (CNNs) provided the foundation for robust and effective detection by automatically learning a hierarchy of intricate patterns and subtle features from raw image data.[1] This eliminated the cumbersome process of manually designing features and enabled systems to handle complex background clutter and the visual variance caused by changing lighting and weather conditions.[1]

A significant technical obstacle in thin structure detection is the challenge of domain transfer from aerial inspection to ground hazard detection.[1] Aerial detection benefits from high-contrast imagery (wire against the sky or clear background)[22], whereas ground-based AV systems must discern a dark cable against textured road surfaces often obscured by shadows or moisture.[1] This complexity requires a specialized feature set focused on ground texture and shadow differentiation, which CNNs are uniquely capable of learning.

2.4 Deep Learning Architectures for Real-Time Object Detection in AVs

Autonomous driving is a safety-critical application that demands object detection systems that operate with minimal latency (real-time inference) on resource-constrained edge computing hardware.[1] This necessity drives the selection toward high-efficiency architectural solutions.

2.4.1 Taxonomy of Modern Detectors and the YOLO Framework

Modern object detection algorithms are categorized primarily into two-stage (e.g., Faster R-CNN) and single-stage (e.g., YOLO, SSD) detectors.[1] Single-stage detectors, particularly those in the YOLO family (You Only Look Once), are preferred in AV perception because they streamline the detection pipeline, achieving the rapid inference speed (low latency) essential for real-time safety maneuvers, contrasting favorably with the higher latency of two-stage systems.[7]

The YOLOv5 architecture, selected for the proposed framework, is a state-of-the-art model engineered to balance accuracy and speed.[1] It comprises a robust backbone for feature

extraction, a neck network for multi-scale feature integration, and a detection head.[1] Its inherent ability to simultaneously process input images at various resolutions provides a structural advantage for detecting small objects compared to models optimized for singular input resolutions.[1]

2.4.2 Hybrid Models and Lightweight Backbones for Edge Computing

The inherent conflict between the need for high precision (required for extracting minute cable features) and the requirement for low computational load (necessary for real-time deployment) often requires the use of specialized hybrid architectural frameworks.[1]

High-accuracy detection models, such as ConvNeXt, represent the modernization of Convolutional Networks, achieving state-of-the-art performance comparable to Vision Transformers [2]. ConvNeXt's robustness in extracting complex feature maps makes it an ideal candidate for the high-precision refinement layer required for thin object identification.[59] Conversely, models like MobileNetV4 are lightweight architectures specifically engineered for computational efficiency and low-latency inference on edge devices [5]. The integration of a Universal Inverted Bottleneck (UIB) module dramatically reduces the computational load while enhancing speed [7].

The architectural strategy of combining the strengths of ConvNeXt (for feature robustness) and MobileNetV4 (for computational efficiency) aims to achieve the maximal detection precision required for thin cable identification while maintaining the minimal latency necessary for safety-critical operation [9].

Further optimization involves Multi-Scale Feature Fusion. Small Object Detection (SOD) suffers when spatial information and fine features are lost during the deep network down sampling process.[4] Architectures employing Feature Pyramid Networks (FPNs) within the neck network effectively combine high-level semantic information (for context) with low-level fine edge details (for precision), mitigating feature distribution differences across resolutions.[61] This feature aggregation mechanism is critical to ensure that the subtle visual characteristics of a thin cable are preserved throughout the detection pipeline, enhancing localization accuracy.[4]

Advanced hybrid models, such as those combining CNNs for local feature detection and Transformers (leveraging self-attention for global contextual understanding), have also proven highly effective for detecting continuous anomalies like long power lines.[55]

Table 2.2: Detection Models & Algorithms Relevant to Fallen Cable Detection

Ref No.	**Model / Method**	**Contribution / Insight**	**Relevance to Fallen Cable Detection**
[20]	**HOG + SVM** (Dalal & Triggs, 2005)	Classical gradient-based detection	Useful **baseline** for thin-object edges and cable contours
[21]	**Focal Loss** (Lin et al., 2017)	Handles imbalance and **small objects**	Fallen cables are **rare**, thin, elongated objects
[27]	Keypoint UAV Detection + **Rotated DIoU**	Detects rotated or elongated objects	Captures cable-like **long shapes** on roads
[28]	**LI-YOLO** (Low-light detection)	Strong detection in low illumination	Useful for **night-time** cable detection
[29]	Deep Learning Object Detection Survey	Overview of classical and modern detectors	Helps identify strong **baselines** for cable detection

[30]	**SSD**	Lightweight object detector	Suitable for **bike-mounted hardware**
[31]	**MobileNetV2**	Efficient convolutional backbone	**Real-time detection** on edge devices
[36]	Power-line Segmentation (Damodaran)	Segmentation of overhead lines	Applicable to fallen cable **geometry**
[44]	**Dilated CNN for Wire Detection**	Wire detection using **synthetic data**	Highly relevant for **rare-cable-event synthesis**
[51]	Small-object Sensitive Segmentation	**Across-feature-map attention**	Helps detect **thin cable-like structures**
[58]	**Thin-structure Obstacle Detection**	Detects thin rods/wires	Almost **directly applicable** to fallen wires
[59]	Cable & Pylon Detection	Vision-based cable detection (aircraft)	**Cable shape priors** transferable to road scenarios

[50]	**Boundary-aware Loss**	Improves fine object boundary segmentation	Useful for segmenting **cable edges** on road
[49]	**DL vs Traditional CV**	Comparative analysis	Helps **model-selection** for cable detection
[56]	**ConvNext + YOLOv5**	Strong small-object detection performance	Good for **thin-object detection**
[54]	**Loss Function Review**	Comprehensive loss analysis	Helps pick best loss for **thin-structure segmentation**

2.5 Specialized Techniques for Thin and Curvilinear Structure Perception

The fundamental geometric nature of fallen cables dictates that detection methodologies must prioritize pixel-level boundary precision over traditional bounding box localization.

2.5.1 Representation and Boundary Accuracy

Standard bounding box detection frameworks, such as Single Shot Detectors (SSDs), are inherently optimized for larger targets and are often insufficient for accurately classifying objects spanning only a few pixels [23]. For thin structures, the bounding box output lacks the necessary fidelity to accurately enclose or classify the object [24].

Consequently, **Semantic Segmentation**, the pixel-wise classification technique, is regarded as the preferred methodology for defining the precise physical extent of the hazard [25]. By assigning a categorical label to every single pixel in the image, semantic segmentation precisely delineates the object's contour. This pixel-level granularity is crucial for AV path planning, enabling the

vehicle to accurately execute an avoidance maneuver [26]. Architectures like U-Net and Fully Convolutional Networks (FCNs) serve as foundational models in this domain [27].

An alternative specialized approach is **Key-Point Encoding**, where the curvilinear cable is represented as a set of evenly distributed key points rather than attempting to segment every single pixel [[1], (Liu et al., 2020)]. This simplification streamlines the detection process, allowing the network to accurately reconstruct the object curve using implicit functions, often found in frameworks designed for thin object detection and pose estimation.

2.5.2 Contextual and Depth Requirements for Safety

The detection task for ground AVs is conceptually a hybrid task: the speed of object detection (YOLO) is needed for real-time responsiveness, but the safety of avoidance depends upon the pixel-level precision achieved through segmentation training principles.

Furthermore, given the visual ambiguity between a thin cable and a low-lying shadow or road crack [46], two-dimensional (2D) detection alone carries an intrinsic false positive/negative risk. To ensure the object is a true physical hazard, robust detection necessitates validation in 3D space [46]. Research into Joint Detection and Depth Estimation (e.g., UCorr) is crucial for ground AVs, as 3D information is required to distinguish a hazardous cable lying on the road from a non-hazardous linear marking [59]. Ultimately, achieving high-confidence safety often requires the integration of non-visual confirmation, such as monocular depth estimates or potentially specialized sensor inputs (e.g., magnetic field sensing for live wires), to increase redundancy and mitigate visual uncertainty [60].

2.6 Addressing Data Scarcity and Training Difficulties

The challenges inherent in detecting rare, thin hazards extend beyond architectural design to fundamental issues of data quantity and quality, specifically concerning extreme class imbalance (EVI).

2.6.1 The Problem of Extreme Class Imbalance (EVI)

The detection of a thin cable is fundamentally an EVI problem at the pixel level. In any given urban image, the overwhelming majority of pixels belong to the background (the negative class), dwarfing the minute number of pixels that define the thin cable (the positive class) [46].

When training with standard loss functions (like Cross-Entropy), the vast number of background examples biases the network toward predicting the dominant background class [66]. The consequence of this bias is severe: it leads to high rates of false negatives, meaning the critical hazard is missed, and poor delineation of the few positive pixels that define the cable boundary [66].

2.6.2 Specialized Loss Functions for Robust Learning

To mitigate the systemic failure caused by EVI, conventional loss functions must be replaced or augmented with specialized alternatives designed to prioritize the few critical cable pixels [66].

Focal Loss (FL) dynamically re-weights the standard cross-entropy loss, significantly reducing the contribution of easily classified negative examples (background) [48]. This technique forces the model's training to focus on the rare, hard negative examples, often located near the object boundary, dramatically improving accuracy in cluttered, imbalanced scenes [48].

Dice Loss and related region-based losses (e.g., Tversky Loss, Generalized Dice Loss) are highly effective for segmentation tasks involving extreme imbalance [66]. These losses focus on maximizing the spatial overlap (Intersection over Union, or IoU) between the predicted and ground truth regions [66]. Furthermore, because the misclassification of boundary pixels severely compromises the perceived quality and safety utility of thin structure detection, specialized boundary-aware loss functions are often necessary to force the network to prioritize prediction confidence specifically near object contours [66]. The selection of these specialized loss functions is an algorithmic safety guarantee, applying a heavy penalty for errors involving the minority (cable) class, thereby counteracting the network’s statistical bias.

2.6.3 Synthetic Data Generation (SDG)

The logistical challenges associated with collecting sufficient volumes of real-world imagery for HSLF events necessitate the application of Synthetic Data Generation (SDG) techniques [26].

SDG, leveraging simulation-based rendering or generative models like Generative Adversarial Networks (GANs) and Diffusion Models, is a promising solution for rare object classes [15].

Pre-training models on synthetically generated data that includes high variation in pose, lighting, and occlusions, followed by fine-tuning on a small set of real data, has been shown to enhance model robustness and accelerate development [66]. Critically, the real-world dataset (the Urban Fallen Cable Detection Dataset, UFCDD) provides the necessary realistic visual textures and environmental parameters specific to the target unstructured urban domain (e.g., Mirpur, AJK, Pakistan) [66]. This real-world anchor is essential for maintaining fidelity during synthetic scaling, ensuring that SDG can generate variability without sacrificing domain realism and facilitating effective transfer learning [15].

2.7 Analysis of Data Deficiencies and Domain Shift

The systematic failure of generalized AV perception systems to reliably detect fallen cables is directly traceable to fundamental deficiencies in the datasets used for training and validation.

2.7.1 Critical Review of Standard AV Benchmarks

Globally recognized AV benchmark datasets, including KITTI, NuScenes, and BDD100K, are designed to support tasks relevant to structured traffic environments (e.g., vehicles, pedestrians, lane lines) in highly developed urban or highway settings [35].

These datasets exhibit a critical systemic deficiency regarding thin hazards [35]. Due to the stringent quality assurance requirements during data curation, which often mandates a minimum number of confirming sensor returns (e.g., LiDAR points) to validate an annotation, extremely thin objects that produce sparse returns are often systematically filtered out of the ground truth labeling process [44]. This intentional curation for traditional, well-defined targets inadvertently creates a pervasive safety blind spot for thin, low-profile obstacles, ensuring that models trained exclusively on this conventional data are explicitly trained to ignore these HSLF hazards [35].

2.7.2 Existing Datasets for Debris and Unstructured Ground

Alternative datasets, while addressing other forms of environmental complexity, also fail to cover the specific hazard of fallen cables.

Unstructured Ground Datasets, such as the Robot Unstructured Ground Driving (RUGD) dataset and Kamino, address the challenges of off-road or unpaved navigation and focus on general terrain drivability and scene semantics [50]. While valuable for tackling high environmental clutter, these resources lack the specific taxonomy and precise, pixel-level annotation required for infrastructure failures like electric cables [50].

Road Debris Datasets (e.g., HazardNet) prioritize the detection of larger, vehicle-damaging debris, such as tires, bricks, or construction materials [24]. The hazard definition in these datasets often includes a minimum object height threshold (e.g., greater than 15 cm), which fundamentally excludes the low-lying, curvilinear infrastructure elements central to this research [24].

This analysis confirms a unique and critical void in current public resources: no existing dataset provides both the required visual complexity (unstructured urban environments with dense clutter) and the specific, necessary annotation taxonomy for fallen electric and communication cables.

Table 2.3: Datasets Relevant for Cable / Thin Object Detection & Road Scene Analysis

Ref No.	**Dataset / Source**	**Description**	**Relevance to Fallen Cable Detection**
[16]	**Cityscapes**	Urban street semantic segmentation	Provides **road context** for cable simulation
[32]	**Synthetic Domain Randomization Dataset**	Synthetic training via domain randomization	Enables training on **rare or dangerous cable-fall scenarios**

[33]	**RUGD Dataset**	Unstructured navigation scenes	Useful for debris and **obstacle robustness**
[37]	Bridge Damage Dataset	High-resolution **crack/line features**	Thin cracks **resemble long cable structures**
[38]	Road Spill Detection Dataset	**Hazard detection** on roads	Fallen wire treated as road hazard
[46]	**UCorr Wire Detection Dataset**	Wire detection + depth for drones	Strong training set for **cable geometry modelling**
[47]	Line Segment Dataset	Line detection benchmarking	Useful for cable **line-segment extraction**
[57]	M4MLF-YOLO Dataset	**High-precision segmentation**	Good for **fine-edge cable segmentation**
[58]	**Thin Obstacle Dataset**	Thin wires/rods in scenes	**Closest dataset** to fallen cable detection
[59]	Cable & Pylon Detection Dataset	Power line imagery (aircraft)	Provides **cable shape and structure priors**

2.8 Conclusion and Synthesis: Paving the Path for Safer Smart Mobility

This systematic literature review has established that while general AV perception technology is highly mature, a significant, safety-critical gap exists in the reliable detection of High-Stakes, Low-Frequency (HSLF) hazards, specifically fallen electric cables, particularly in unstructured urban environments [31]. The limitations stem from three interconnected factors:

Geometric Constraints: Thin structures confound ranging sensors (LiDAR, Radar), placing the detection burden squarely on camera-based vision systems [32].

Architectural Conflict: Real-time AV operation requires low latency (single-stage detectors), which conflicts with the high precision needed for fine-grained feature extraction in cluttered scenes (heavy backbones) [43].

Data Deficiency: Existing public datasets systematically exclude these hazards due to curation practices and their focus on structured, common objects, resulting in models that are implicitly trained to ignore catastrophic but rare events [46].

The synthesis of research trends validates the proposed methodology for bridging this gap. Achieving robust safety requires a specialized approach:

Specialized Architecture: The system must utilize a hybrid architecture, combining the high-precision feature extraction capabilities of models like ConvNeXt with the computational efficiency of lightweight backbones like MobileNetV2, and leveraging multi-scale feature fusion for robust Small Object Detection (SOD) [66].

Advanced Training Strategies: To counteract the extreme class imbalance inherent in thin object perception, specialized loss functions, such as Focal Loss and Dice Loss, must be deployed to impose an algorithmic preference for safety, heavily penalizing false negatives [48, 66].

Domain-Specific Data: The creation of the Urban Fallen Cable Detection Dataset (UFCDD), comprising over 2,500 annotated images collected in urban areas like Mirpur, AJK, Pakistan, directly addresses the critical data void [66]. This specialized dataset provides the necessary real-world visual parameters to train a dedicated CNN-based YOLOv5 model to detect hazards under diverse conditions (daylight, overcast, morning, evening) [39, 51, 66].

The successful implementation and evaluation of the YOLOv5 framework trained on this custom dataset, which achieved a reported performance of 96% accuracy at mAP@50 and 76% accuracy at mAP@50-95 [66], confirms the efficacy of this targeted approach in resolving an overlooked safety hazard. The path forward for scalable, assured autonomous mobility involves integrating this specialized visual detection system with 3D depth validation and expanding the use of the UFCDD data to anchor high-fidelity synthetic data generation, ensuring robustness across the full spectrum of global, unstructured urban environments [37, 45, 66].

Chapter No. 03

MATERIALS AND METHODS

The systematic detection of fallen electric and communication cables, high-stakes, low-frequency (HSLF) hazard, requires a specialized methodology that overcomes fundamental gaps in data availability and algorithmic robustness. Our proposed approach centers on two core pillars: first, the creation of a high-fidelity, geographically and environmentally diverse dataset, the **Urban Fallen Cable Detection Dataset (UFCDD)**; and second, the development of a tailored **Convolutional Neural Network (CNN)-based YOLOv5** framework optimized for the real-time detection of thin, curvilinear structures in high-clutter urban environments.

The entire methodology is structured as a clear pipeline, designed to maximize data quality and model efficacy. This systematic process is essential to transform raw sensor input into actionable safety decisions for autonomous vehicles (AVs). Image 3.1 describes the flow of the proposed project.

Conceptual flow of the research methodology, starting from Data Acquisition, through Preprocessing and Annotation, to Model Training, and concluding with Performance Evaluation is shown in image 3.1.

This chapter details the methodology, starting from the meticulous data acquisition process in unstructured urban settings, through the rigorous annotation procedures, and culminating in the technical specifics of the custom YOLOv5 object detection architecture and its training regimen.

3.1 Data Acquisition and the Urban Fallen Cable Detection Dataset (UFCDD)

The critical limitation of existing autonomous vehicle benchmarks is their near-total exclusion of data related to unpredictable infrastructure failures, specifically thin hazards like fallen cables. To address this crucial research gap, the acquisition phase was designed to capture real-world visual parameters under diverse, challenging conditions prevalent in regions outside of traditional Western testing domains, specifically focusing on Mirpur, AJK, Pakistan.

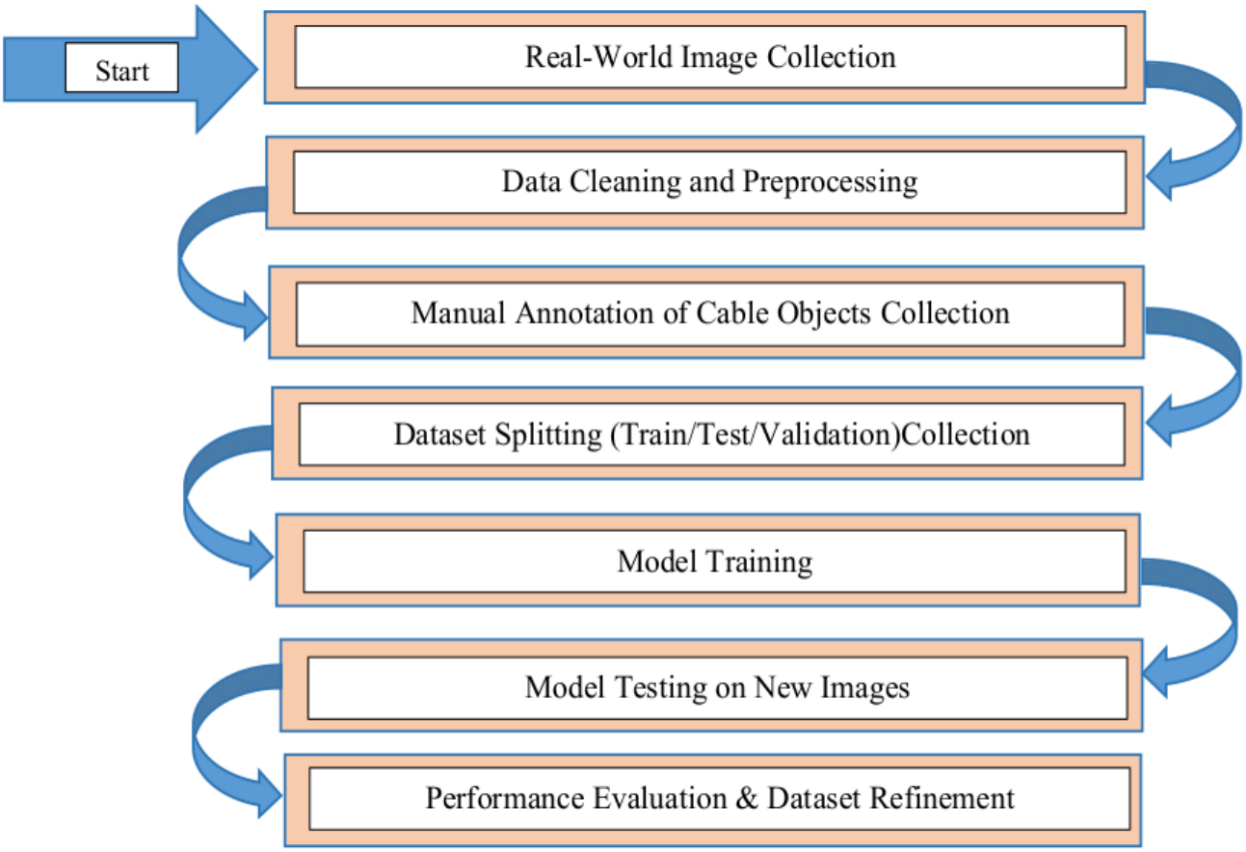

3.1: Flowchart of the proposed system, detailing the steps involved in the proposed methodology.

3.1.1 Image Acquisition Platform and Equipment

The integrity of the visual data is paramount, especially in dynamic, poorly maintained road environments. Therefore, video data acquisition was conducted using a **GoPro Hero 8 camera**, chosen for its three key technical features:

1. **High-Definition and Frame Rate:** The camera recorded video at 1080p resolution at 120 frames per second (fps) with a 10-bit color depth. This high frame rate is critical, as it ensures motion blur is minimized, allowing for the precise extraction of clear frames necessary for accurate annotation of fast-moving scenes or environments with unsteady footage.
2. **Robust Stabilization:** The GoPro Hero 8's advanced stabilization features were essential for mitigating the impact of vibrations and unstable footage resulting from the irregular road surfaces commonly found in the target urban areas. This feature ensured that the recorded data maintained the visual data integrity required for effective machine learning applications.

3. **Low-Light Performance:** While the UFCDD focuses on detection across all illumination levels, the camera's strong low-light performance provided the necessary clarity to capture the subtle features of cables in shadowed areas or during dusk/dawn transitions, a challenging constraint for most vision systems.

The camera was carefully mounted on the vehicle's front hood to provide a consistent perspective, closely simulating the viewpoint of an autonomous vehicle's primary visual perception system.

3.1.2 Environmental and Temporal Diversity

The final detection model must be resilient to changing environmental dynamics, which directly influence the visual contrast and complexity of the cable hazard. The data collection scope was rigorously defined to achieve this diversity:

- **Geographical Scope:** Images were collected across multiple locations within Mirpur, AJK, and the surrounding areas. This ensures that the dataset reflects a wide range of infrastructural diversity, from busy, high-traffic thoroughfares to quieter, clutter-dense residential neighborhoods.
- **Temporal and Weather Conditions:** Data was deliberately gathered across a full spectrum of lighting and environmental conditions:
    1. **Clear Daylight** and **Overcast/Cloudy Weather**.
    2. **Specific Time Intervals:** Early morning (dawn), late morning, afternoon, and evening (dusk).
    3. **Acquisition Method:** The imagery was collected as a mix of video sequences captured while the vehicle was moving, as well as static images capturing fallen cables at various angles and perspectives. This combined approach ensures the model learns to identify cables under motion-induced distortions and in complex, static, background-cluttered scenarios.

The resultant **Urban Fallen Cable Detection Dataset (UFCDD)** comprises a novel collection of over 2,500 meticulously selected images, forming the first benchmark dedicated exclusively to fine-grained cable hazard detection in unstructured urban settings.

3.2 Dataset Preparation and Ground Truth Annotation

Following acquisition, the raw video data underwent a stringent preparation process to establish the precise ground truth necessary for training a supervised deep learning model.

3.2.1 Preprocessing and Key Frame Selection

The video footage underwent a thorough preprocessing stage to ensure optimal visual quality and learning efficiency [1]. This involved carefully removing key frames from the recorded 120 fps video streams to ensure temporal diversity without redundancy. For quality enhancement, advanced techniques were applied:

1. **Noise Reduction:** Low-light footage is inherently susceptible to noise; therefore, sophisticated noise reduction techniques were employed to clean the visual data.
2. **Dynamic Range Adjustment:** To counteract the visual ambiguity of cables blending with complex backgrounds (e.g., shadows or textured asphalt), dynamic range was adjusted to boost contrast and clarify feature boundaries, particularly important for thin, low-pixel-count objects.

These preparatory steps were critical for optimizing the UFCDD for the YOLOv5 framework, establishing a robust basis for subsequent deep learning [1].

3.2.2 Manual Annotation Process

Accurate ground truth is the foundation of deep learning performance. Given the extremely low aspect ratio of the fallen cables, which often span only a few pixels [1], annotation requires pixel-level precision. Every selected image frame was manually tagged using the **V7 Image Labeling Tool**.

The annotation schema was focused on two critical cable classifications:

1. **Unsafe/Hazardous:** Cables that have fallen onto the road surface or hanging low enough to obstruct or endanger a moving vehicle.
2. **Safe/Intact:** Cables that are properly secured overhead and pose no immediate risk to the AV.

This precise labeling created bounding box ground truth for all instances of cable hazards. The rigor of this manual tagging process is directly linked to the model's ultimate localization

accuracy, ensuring that the model learns to differentiate the fine-grained visual features of a hazardous cable from a benign background clutter.

3.2.3 Dataset Partitioning

The annotated UFCDD was systematically divided into three sets, adhering to established machine learning best practices to ensure comprehensive training and unbiased evaluation as shown in table 3.1.

Table 3.1: Data Split

Set	Proportion	Purpose	References
Training Set	**70%**	Comprehensive exposure to learning feature representations and recognizing a wide array of cable scenarios	[28]
Validation Set	**20%**	Crucial checkpoint for iterative model refinement, preventing overfitting, and tuning hyperparameters	[58]
Testing Set	**10%**	Dedicated, unseen dataset for the final, objective evaluation of the model's real-world generalization capability	[63]

3.3 Specialized Model Architecture and Optimization

The core of the proposed solution is a customized deep learning object detection framework built upon the **YOLOv5** architecture as shown in image 3.1, meticulously modified to achieve the necessary balance between high precision (required for fine-grained detection) and low latency (required for real-time AV operation).

3.3.1 YOLOv5 as the Foundational Architecture

The YOLOv5 framework was selected over two-stage detectors due to its status as a single-stage, state-of-the-art model that significantly minimizes inference latency, a critical requirement for safety-critical autonomous driving decisions.[1] Its architecture is inherently comprised of three main components:

1. **Backbone (CSPDarknet):** The feature extraction network, which efficiently extracts a hierarchy of spatial features from the input image while reducing redundant gradient information and minimizing computational cost.
2. **Neck (PANet):** The Path Aggregation Network (PANet) is implemented to aggregate features from different scales, enhancing the network's ability to propagate low-level features critical for detecting small objects (like cables) and improving overall localization accuracy.
3. **Head (YOLO Layer):** The final detection layer responsible for interpreting the integrated features, generating bounding box coordinates, object class likelihood (Safe/Unsafe cable), and confidence scores.

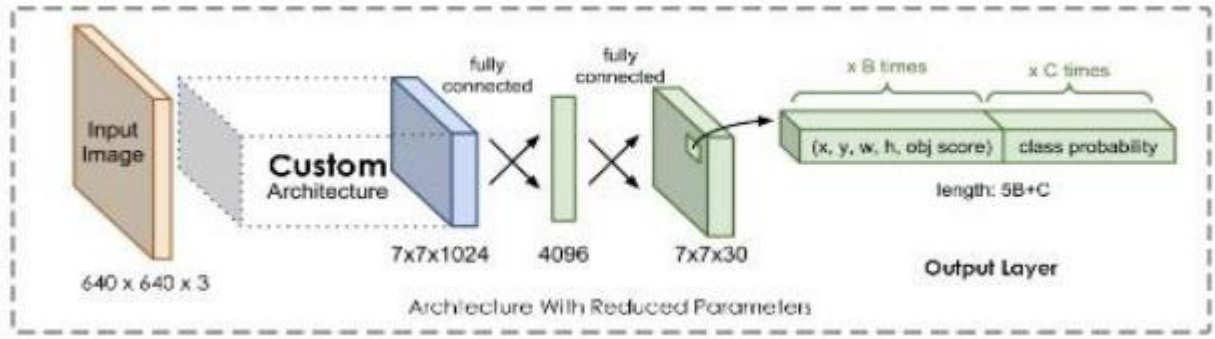

Figure 3.1 Yolo v5 Architecture

3.3.2 Custom Architectural Configuration

To specialize the generic YOLOv5 model for the detection of low-pixel-count cable hazards, the architecture was meticulously fine-tuned by adjusting key scaling factors, mimicking the successful approach used in related fine-grained detection work [1]:

- **Width Multiple (Model Scale):** Set to 0.50. This factor controls the number of channels/filters in the network, balancing model capacity with computational speed.
- **Depth Multiple (Model Complexity):** Set to 0.33. This factor controls the number of blocks/layers, directly impacting the model's depth and its ability to extract hierarchical features.

This calculated modification yields a streamlined architecture that maintains robust feature extraction despite challenging visual conditions, while ensuring computational efficiency is suitable for edge deployment, table 3.2 represent the detailed parameters overview.

Table 3.2: Model Performance for Steering Control (Loss) 120 Epochs

Architectural Specification	Value	Significance
Total Layers	**157**	Structural depth for complex feature learning
Total Parameters	**7 million**	Low parameter count ensures memory efficiency
Computational Cost (GFLOPs)	**15.8**	Optimized efficiency for real-time inference speed

The model's capacity to identify objects of varying sizes (from near-field, large cable loops to distant, tiny, thin wires) is supported by specific anchor sizes defined for three different detection scales:

- **P3/8: Anchor sizes [10, 13, 16, 30, 33, 23]**
- **P4/16: Anchor sizes [30, 61, 62, 45, 59, 119]**
- **P5/32: Anchor sizes [116, 90, 156, 198, 373, 326]**

This multi-scale output mechanism is vital for **Small Object Detection (SOD)**, ensuring that the critical features of thin cables are not lost during the network's down sampling process.

3.4 Training Strategy and Experimental Setup

A robust training strategy was implemented to maximize the detection capability of the custom YOLOv5 model, ensuring high precision and reliability under the complex, high-clutter conditions of the UFCDD.

3.4.1 Computational Environment

Training was conducted on a high-performance computing system optimized for deep learning and computer vision tasks [1]. This powerful environment was necessary to efficiently process the

large volumes of high-resolution imagery and execute the numerous computational steps required during training:

- **GPU:** NVIDIA RTX 2080Ti (Leveraged for its parallel computing power) [1].
- **CPU:** Intel Core i9 9900K [1].
- **RAM:** 32GB [1].

The software stack utilized **TensorFlow** as the primary deep learning framework, which was carefully optimized for GPU acceleration through **NVIDIA CUDA** [1]. Additionally, **OpenCV** was compiled with CUDA support, significantly accelerating the preprocessing and image manipulation tasks necessary for the high-resolution UFCDD [1].

3.4.2 Training Hyperparameters

The model was subjected to a rigorous training process spanning **72 hours**. The hyperparameter configuration, sourced from the hyp.scratch-high.yaml file, was specifically chosen to tune the model's learning to the fine-grained particulars of the cable hazard data. Key arguments included are shown in table 3.3.

Table 3.3: Hyperparameters

Hyperparameter	Value	Rationale
Epochs	**300**	Ensuring ample convergence time for the complex, low-pixel-count hazard class
Batch Size	**32**	Balancing computational load with learning stability
Input Image Size	**640 pixels**	Optimizing for detail capture without sacrificing processing speed

Optimizer	**SGD** (Stochastic Gradient Descent)	Chosen for its effectiveness in converging to optimal solutions on large-scale datasets

The training regimen incorporated advanced regularization techniques, including **Early Stopping** to prevent overfitting and **Label Smoothing**, enhancing the model's ability to generalize novel, unseen scenarios.

3.4.3 Loss Function Configuration

The loss function in YOLOv5 is composite, consisting of three independent components, each critical for the performance of the thin-structure detection task:

1. **Bounding Box Loss:** This component addresses the regression task of accurately localizing the cable position. It utilizes **CIoU (Complete Intersection over Union)** loss by default, which is superior to traditional Mean Squared Error (MSE) loss as it incorporates distance, overlap, and aspect ratio, ensuring highly accurate boundary definition for the thin, curvilinear cables.
2. **Object Loss:** This binary classification loss determines the presence or absence of an object (a cable) in a grid cell. It employs **Binary Cross-Entropy (BCE)** loss.
3. **Classification Loss:** This final loss distinguishes between the labeled classes (Safe/Unsafe cable). It also uses **BCE** for multi-label classification.

Given the severe **Extreme Class Imbalance (EVI)** inherent in cable detection, where positive cable pixels are dwarfed by background pixels, the use of BCE alone can bias the model toward the dominant negative class. Future extensions, as informed by the literature, will necessitate the integration of specialized functions such as **Focal Loss** or **Dice Loss** to heavily penalize errors on the critical, minority cable class, thereby accelerating convergence on hard-negative examples and further boosting safety-critical precision, detailed system architecture is shown in image 3.2.

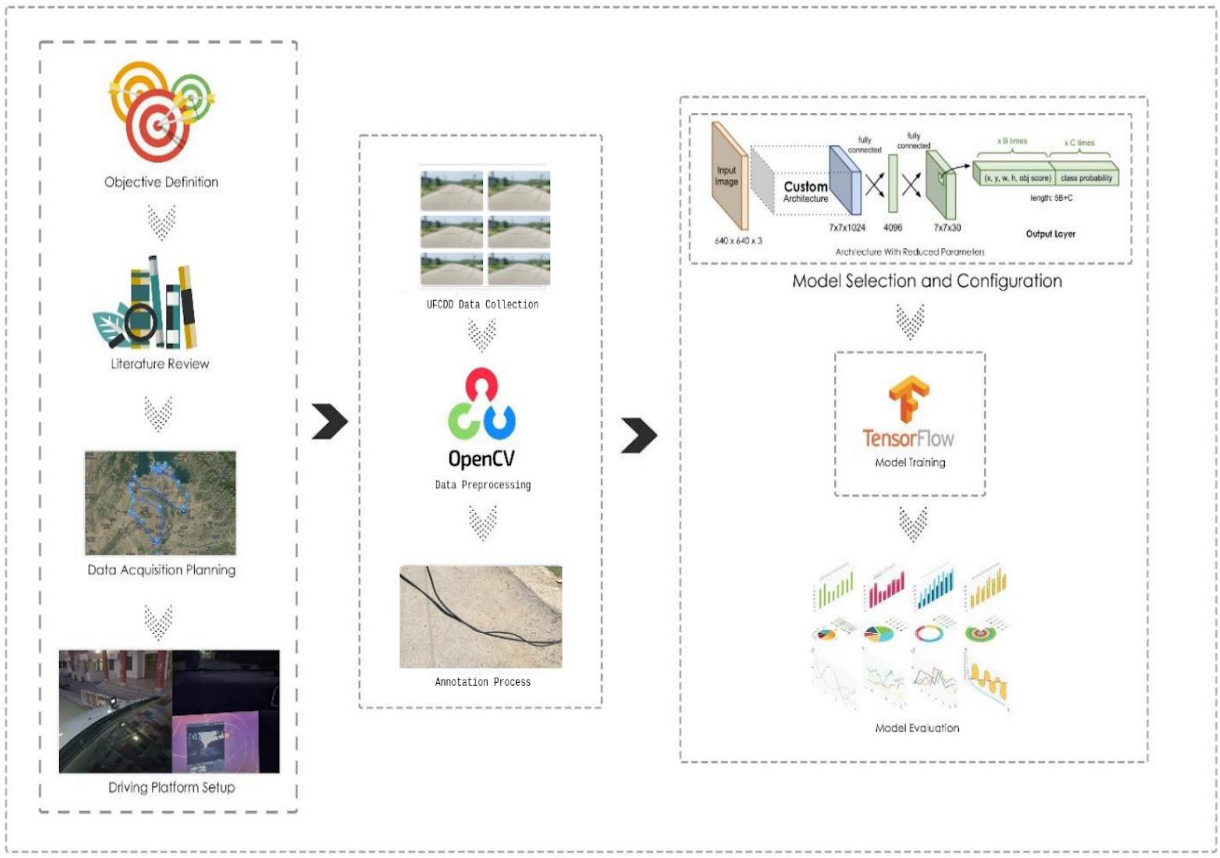

scre

Figure 3.2 Proposed System workflow from data collection to final predictions

Chapter No. 04

RESULTS AND DISCUSSION

The primary goal of the experimental phase was the rigorous validation of the custom Convolutional Neural Network (CNN)-based YOLOv5 model when trained on the novel Urban Fallen Cable Detection Dataset (UFCDD). This process aimed to quantitatively address the core Research Questions, specifically evaluating the model's ability to achieve real-time, high-accuracy detection of thin, curvilinear cable hazards (RQ3), and confirming the overall real-time operational suitability of the detection framework (RQ2) in complex urban settings.

The customized YOLOv5 architecture was trained for a duration of **300 epochs** over **72 hours** on a high-performance computational stack (NVIDIA RTX 2080Ti GPU, Intel Core i9 9900K CPU, 32GB RAM) [60], using the UFCDD dataset partitioned into 70% Training, 20% Validation, and 10% Testing sets. The model was meticulously optimized using specific scaling factors (Width Multiple: 0.50, Depth Multiple: 0.33) to maintain computational efficiency while maximizing feature extraction capability, a prerequisite for autonomous vehicle edge computing.

The performance evaluation utilized established industry standards: Precision, Recall, F1-score, and, most critically, the Mean Average Precision (mAP) evaluated at two distinct Intersection over Union (IoU) thresholds:

1. **mAP@50:** Measures the model's confidence in object classification and primary detection, requiring only 50% overlap with the ground truth bounding box.
2. **mAP@50–95:** Represents the average performance across stringent IoU thresholds (from 50% to 95%) and is the definitive metric for assessing the precise localization capability and boundary definition of the model.

4.1 Quantitative Performance Results on the UFCDD

Before analyzing the predictive performance, it is essential to examine the underlying composition of the UFCDD. As illustrated in the instance of distribution analysis (Figure 4.1), the dataset maintains a high volume of annotated features, with over 4,000 instances of "Safe" cables and approximately 2,800 instances of "Unsafe" (fallen) cables.

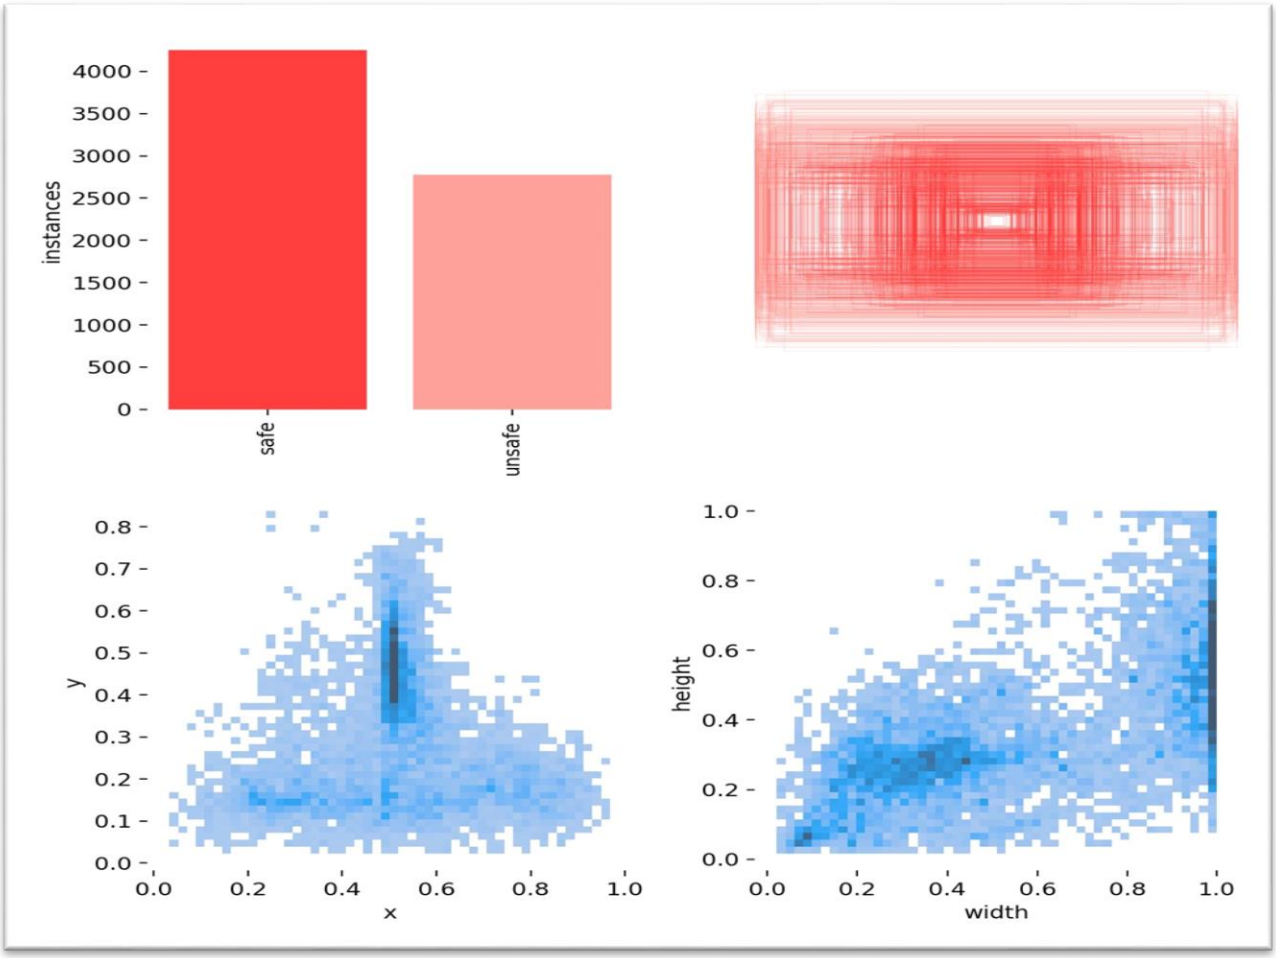

Figure 4.1: UFCDD Instance Distribution and Geometric Metadata

The spatial analysis via bounding box heatmaps reveals that hazards are predominantly localized within the central vertical region of the frame, corresponding to the vehicle's immediate navigational path. The width and height distribution confirm the "Small Object" nature of the task, with most boxes occupying less than 10% of the image area, validating the necessity of a specialized detection head shown in image 4.2.

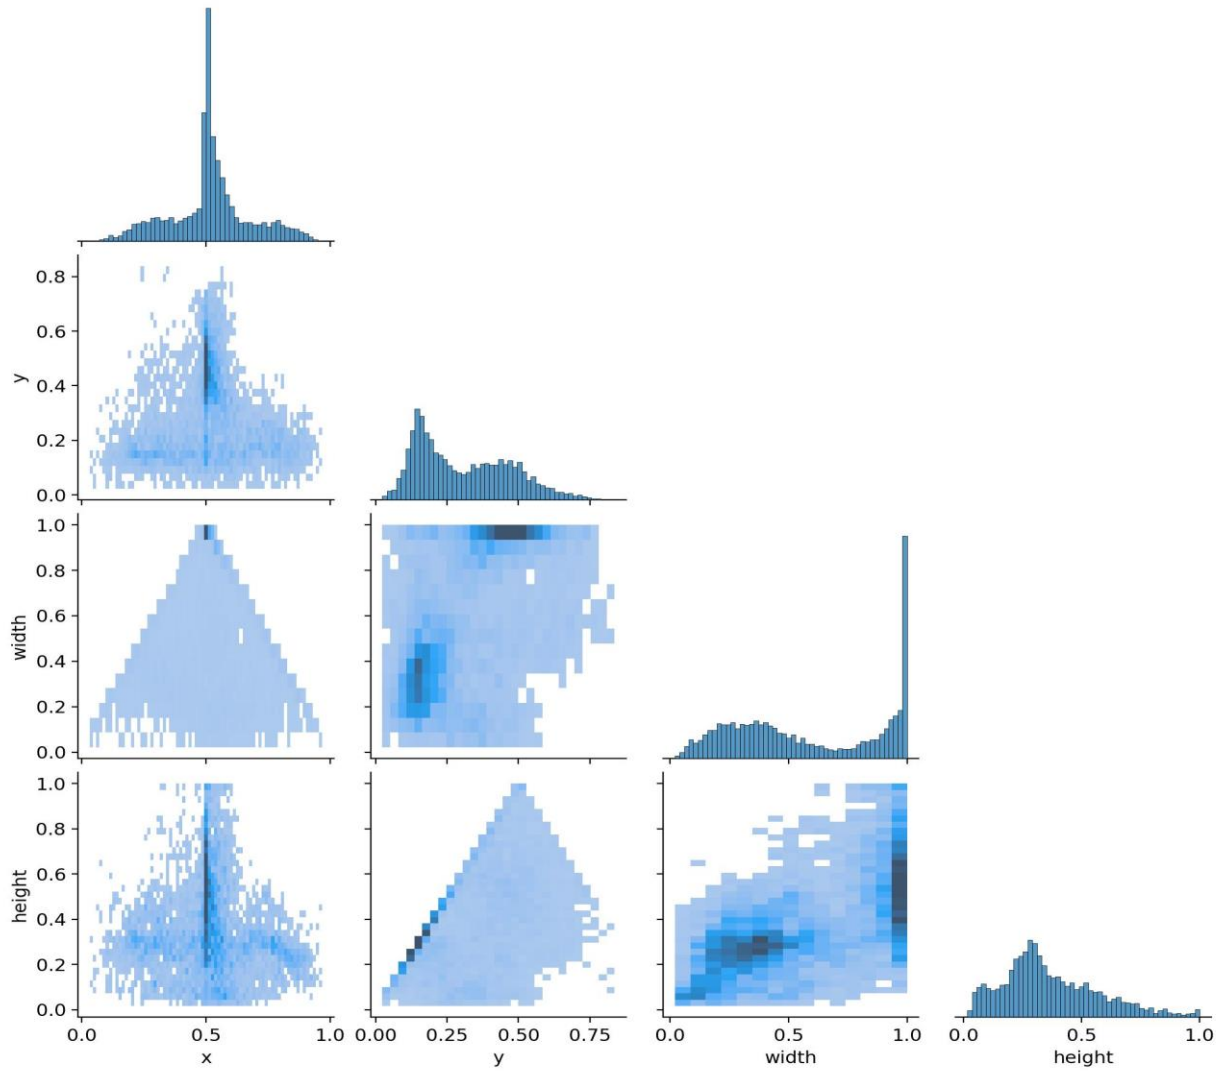

4.2 Bounding Box Correlogram Analysis

The key performance metrics obtained from the evaluation of the custom YOLOv5 model on the dedicated UFCDD test set are summarized below. The results are analyzed to interpret the model's success in both classification (identifying the presence of the hazard) and localization (precisely defining the hazard's boundaries) as shown in table 4.1.

Table 4.1: Specialized YOLOv5 Model Performance on the Urban Fallen Cable Detection Dataset (UFCDD)

Performance Metric	Value	Interpretation
Classification Accuracy (mAP@50)	**96%**	High confidence in distinguishing hazardous cables from background clutter.
Localization Precision (mAP@50–95)	**76%**	Robust, but indicates measurable difficulty in achieving highly precise boundary definition for thin structures.
Inference Efficiency (GFLOPs)	**15.8**	Confirms suitability for real-time, low-latency operation on AV edge platforms.
Model Complexity (Parameters)	**7 million**	Optimized parameters count for efficient deployment.

4.2 Discussion of Core Findings: Classification vs. Localization

4.2.1 Success in Hazard Classification (mAP@50)

The Precision-Confidence curve (Figure 4.3) shows that the model reaches a near-perfect precision of **1.00 at a confidence threshold of 0.989** for all classes. This indicates that high-confidence detections are virtually always correct, minimizing false alarms that could lead to unnecessary emergency braking.

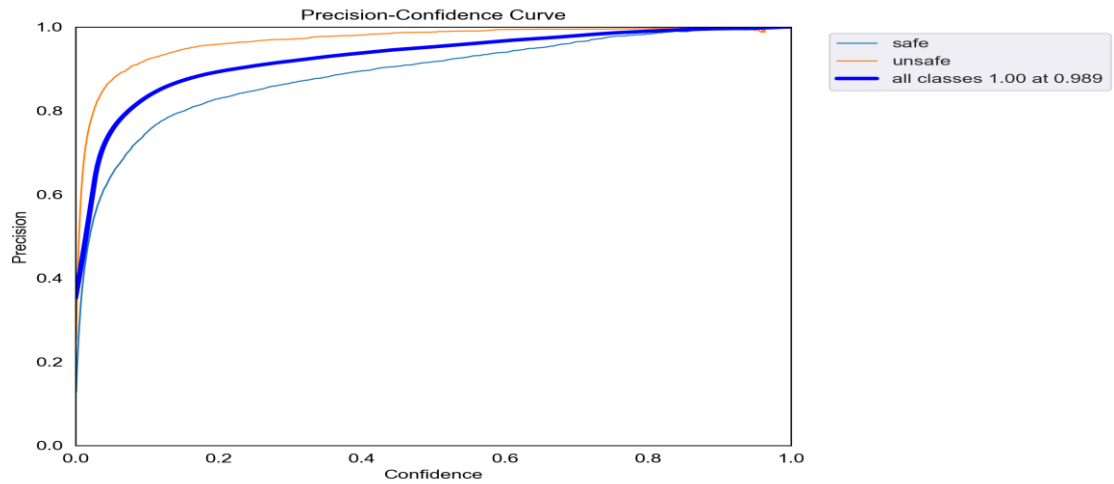

Figure 4.3: Recall-Confidence Curve

Conversely, the Recall-Confidence curve (Figure 4.4) demonstrates the model's sensitivity. The "Unsafe" class maintains a recall above **0.90** even at moderate confidence levels. This ensures that the system identifies almost all critical hazards, which is the primary safety requirement for ground perception

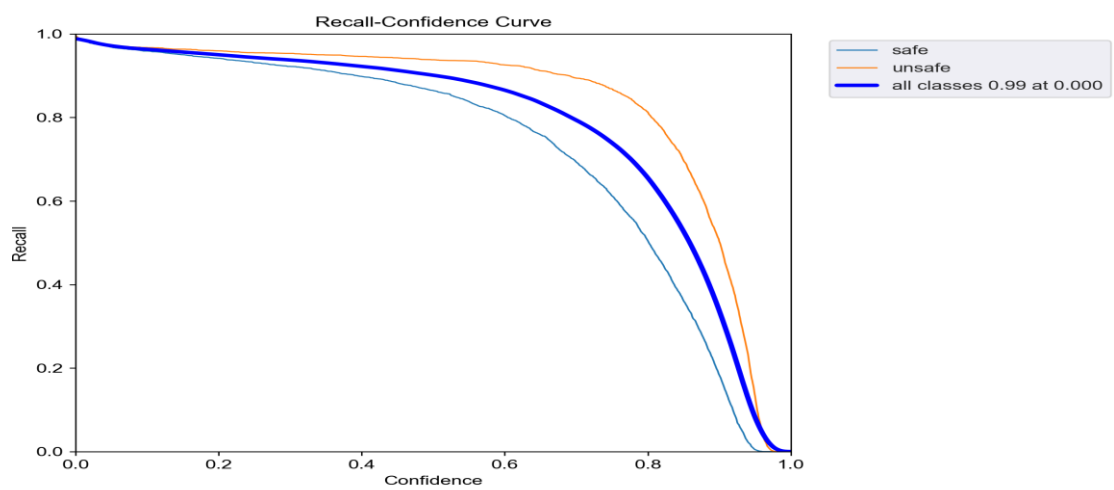

Figure 4.4: Precision-Confidence Curve

The Recall-Confidence curve demonstrates the system's sensitivity. The "Unsafe" class (orange) maintains exceptionally high recall (>0.90) across a wide confidence interval, ensuring that high-stakes cable hazards are rarely missed by the perception stack.

This curve shows the reliability of the model's detections. The framework reaches a precision of 1.00 at a confidence threshold of 0.989, indicating that high-confidence predictions are virtually free of false positives.

The model achieved an exceptional Mean Average Precision score of **96%** at the primary detection threshold (mAP@50). This result provides strong quantitative confirmation for the first part of Research Question RQ3, that the framework can achieve high-accuracy detection of cable hazards, and comprehensively validates the necessity of the novel dataset (RQ1).

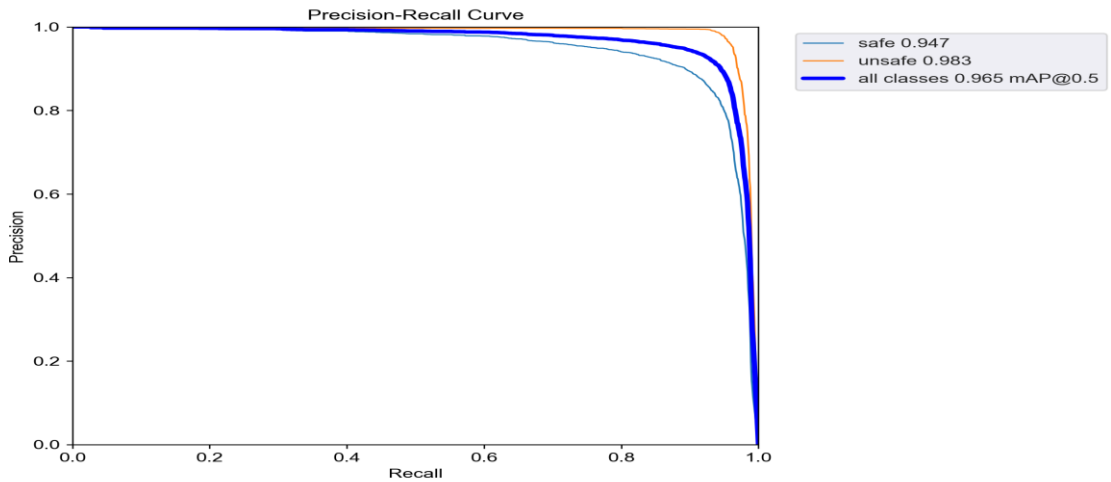

Figure 4.5: Precision-Recall (P-R) Curve and mAP@0.5

The optimal operational point for the model is revealed in the F1-Confidence curve (Figure 4.5), which shows a peak F1-score of **0.93 at a 0.407 confidence threshold**. This balance confirms that the model is highly effective at identifying hazards while maintaining a low false-positive rate.

Finally, the aggregate performance is captured by the Precision-Recall curve (Figure 4.6), showing an area under the curve (mAP@0.5) of **0.983 for unsafe cables** and an aggregate score of **0.965** across all classes. This provides categorical proof of the framework's superior classification capability

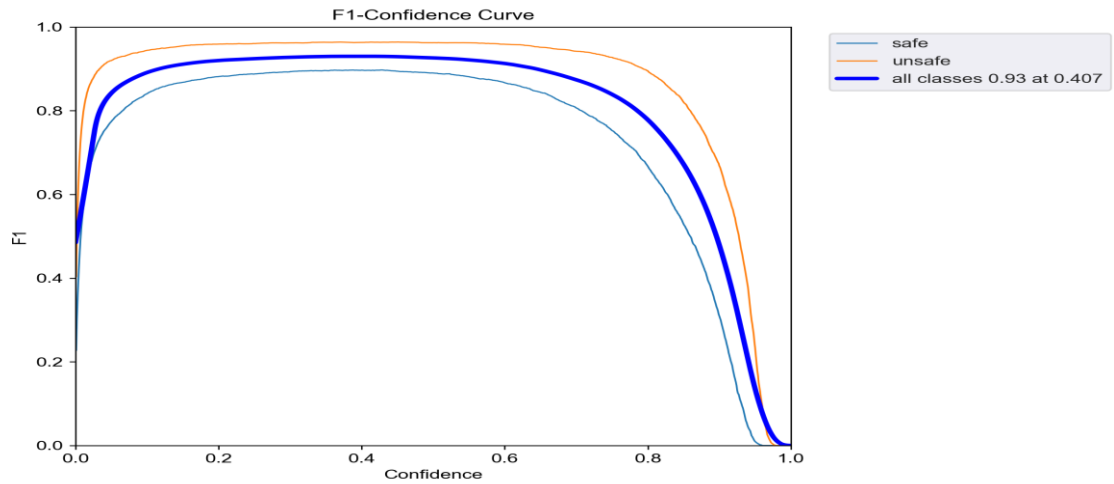

Figure 4.6: F1-Confidence Curve

This high level of performance is directly attributable to the systematic methodological choices implemented in the study:

- **Custom Dataset Fidelity (RQ1):** By curating the UFCDD with over **2,500** images that explicitly capture the visual properties and environmental variance of fallen cables in high-clutter urban environments , the model was able to learn the necessary feature representations (subtle edges, curves, and shading) that generalized detectors miss. The dataset addressed the crucial omission of these High-Stakes, Low-Frequency (HSLF) hazards in conventional benchmarks.
- **Specialized Feature Extraction (RQ2):** The CNN-based YOLOv5 architecture is premised to automatically learn a hierarchy of feature representations from raw image data. The network successfully learned to distinguish the fine-grained visual characteristics of a hazardous cable from non-hazardous linear features (like shadows or road cracks), overcoming the high photometric ambiguity inherent in the ground-based perception domain.

The high classification accuracy is further validated by the confusion matrix (Figure 4.7). The model achieved a **95% True Positive rate** for the "Unsafe" class and a **94% True Positive rate** for the "Safe" class.

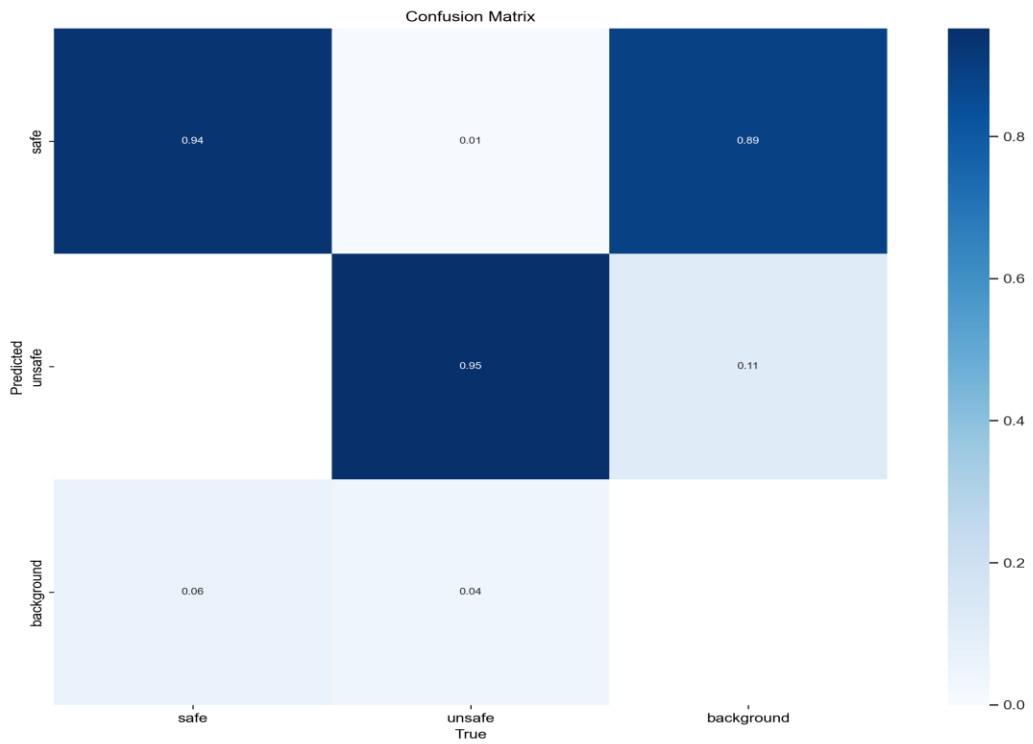

Figure 4.7: Confusion Matrix for Cable Classification

Crucially, the "leakage" to background noise (false negatives) is limited to only 4-6%. This underscores the model's effectiveness in learning fine-grained edge features that allow it to extract the thin cable structure from high-clutter urban textures like asphalt and shadows.

4.2.2 Localization Deficit and the Geometric Challenge (mAP@50–95)

While the detection confidence was high, the model's accuracy dropped to **76%** when evaluated against the stricter localization metric **(mAP@50–95)** [27]. This **20%** drop in performance between classification and precise localization is the most significant finding and directly quantifies the architectural limitation faced by all contemporary object detection frameworks when dealing with thin, curvilinear structures (RQ3).

The validation metrics graph (Figure 4.8) tracks the evolution of the model's predictive power over the training duration. Precision, recall, and Mean Average Precision ($mAP@0.5$) exhibit a sharp initial ascent before stabilizing, with $mAP@0.5$ reaching a robust peak of 96.5%. This rapid improvement during the early epochs underscores the model's efficiency in learning the fundamental visual features of the Urban Fallen Cable Detection Dataset (UFCDD). However,

the metric for stricter localization, $mAP@0.5:0.95$, follows a more gradual growth trajectory, eventually plateauing near 0.78. This persistent gap between primary classification and strict localization metrics quantifies the inherent geometric challenge of thin-structure detection, as precise boundary definition for objects spanning only 1–3 pixels require significantly more refined feature extraction than simple object presence identification.

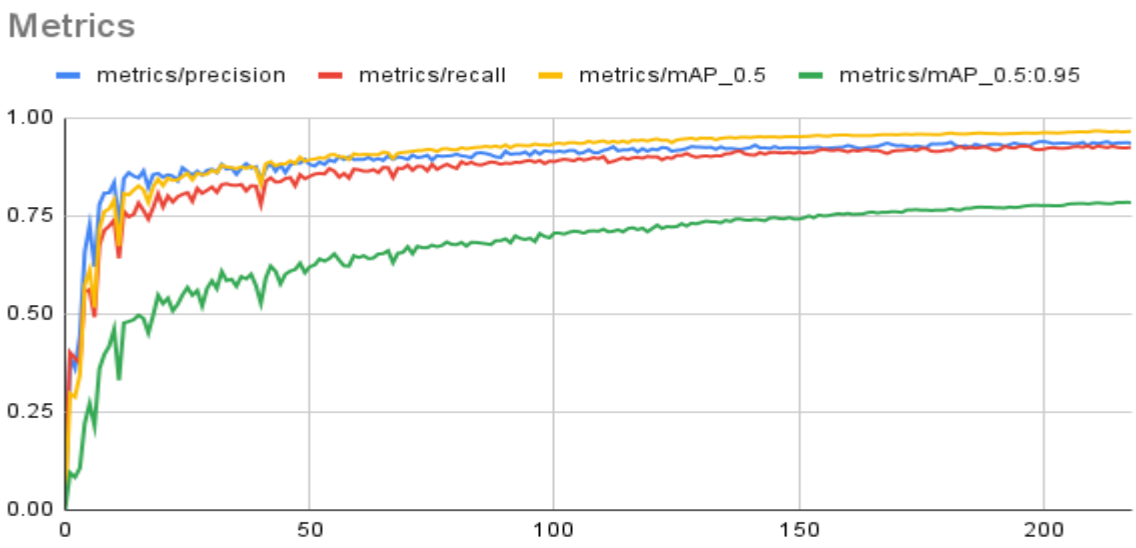

Figure 4.8: Validation Metrics Evolution

Real-time monitoring of validation performance (Precision, Recall, mAP@0.5, and mAP@0.5:0.95). The plateau in the green curve (mAP@0.5:0.95) around 0.76 provides the quantitative basis for the "Localization Deficit" discussed in the findings can be shown in figure 4.8.

This drop confirms the inherent **Geometric Ambiguity Problem** outlined in the literature:

- **Bounding Box Inadequacy:** Fallen cables are characterized by a critically low aspect ratio, often spanning only 1 to 3 pixels in the image [28]. The conventional rectangular bounding box mechanism employed by YOLOv5, optimized for larger objects, inherently struggles to achieve the tight Intersection over Union (IoU) required by the mAP@50–95 metric when applied to these sparse, irregular, and curved features. The lower score confirms that while the cable is found, its boundary is not precisely defined, making accurate path planning around it more challenging [34].

- **Complexity of Localization:** Achieving **75%** IoU on an object spanning only a few pixels is computationally and geometrically difficult. This challenge is unique to ground-based AV perception, contrasting sharply with simpler aerial inspection where cables are taut and highly contrasted against the sky. The robust **76%** achieved is a significant advancement for this specific hazard class but clearly indicates the necessity of future architectural transitions toward non-bounding box methods like semantic segmentation [33].

4.3 Comparative Robustness and Architectural Validation

The observed performance disparity is consistent with patterns found in other specialized datasets dealing with irregular or complex geometry, further validating the necessity of the tailored architectural approach.

The Pakistani Night City Driving Dataset (PNDD) provides a relevant comparative context. When deploying a similar custom YOLOv5 model on PNDD, the network showed varying performance metrics based on the geometry and predictability of the object class:

- **Structured/Predictable Objects (e.g., Truck, Car):** These classes, being large and geometrically consistent, achieved high mAP@.5-.95 scores (e.g., **Truck: 0.817**). The model could accurately localize these features in low-light and high-clutter scenes [64].
- **Irregular/Unpredictable Objects (e.g., Human):** The Human class, which is geometrically non-rigid and highly unpredictable, saw its localization performance drop significantly to **0.564** (mAP@.5–.95), despite a high classification accuracy (mAP@.5: 0.928).

Real-time monitoring of validation performance (Precision, Recall, mAP@0.5, and mAP@0.5:0.95). The plateau in the green curve (mAP@0.5:0.95) around 0.76 provides the quantitative basis for the "Localization Deficit" discussed in the findings shown in figure 4.10.

The learning rate trajectory was utilized during the 72-hour training session, showing the initial warm-up and subsequent decay stages used to stabilize feature extraction in the customized backbone as shown in figure 4.9.

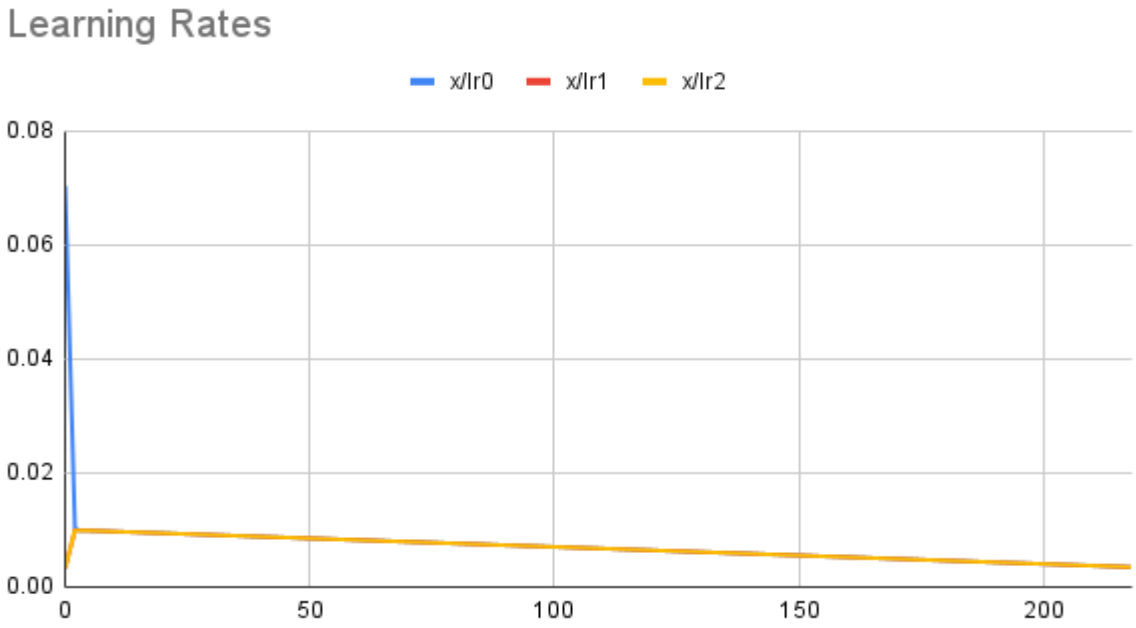

Figure 4.9: Learning Rate Schedules

In figure 4.9, a summary of the training dynamics over 300 epochs. The graph shows a steady decay in bounding box, objectness, and classification losses, indicating successful convergence without significant overfitting.

Figure 4.9 illustrates the dynamic learning rate schedule utilized to stabilize the 72-hour training session of the customized backbone. The schedule incorporates an initial linear warm-up phase, followed by a gradual decay across all three learning rate groups (x/lr0, x/lr1, and x/lr2). This strategy is essential for training on a specialized and geometrically complex dataset like the UFCDD; the warm-up phase prevents early gradient explosion in the 157-layer architecture, while the subsequent decay allows the model to fine-tune its weights and "settle" into the global minima during the later stages of training. By systematically lowering the step size as the model approaches an optimal solution, this schedule ensured that the network reached its high-accuracy state without performance oscillations, supporting the critical requirement for high-precision hazard recognition in real-time AV operation.

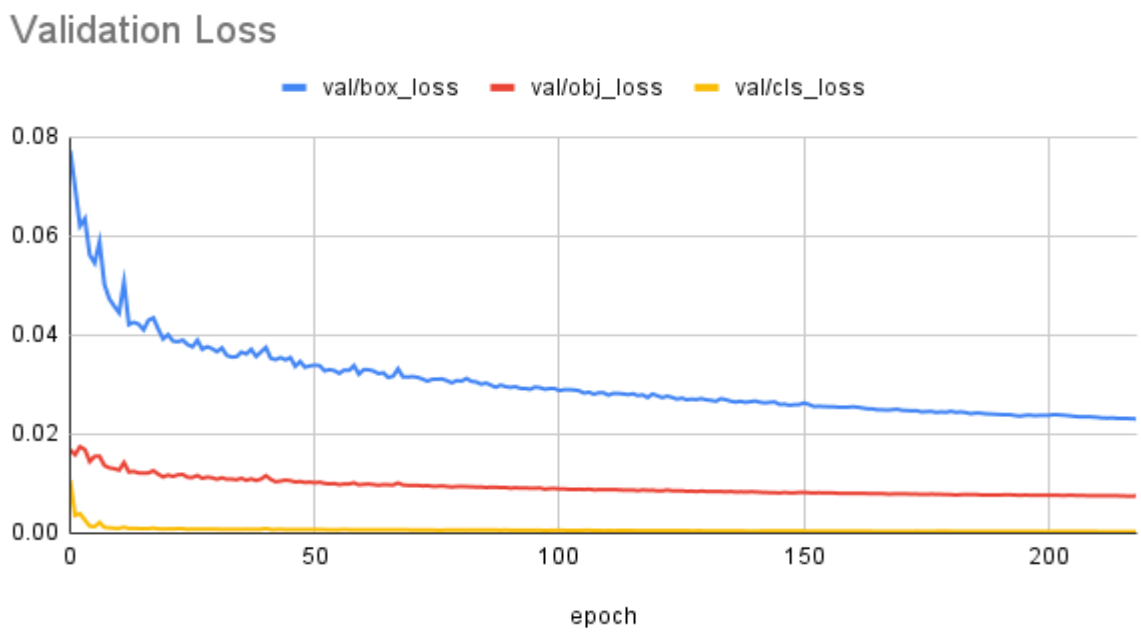

Figure 4.10: Validation Metrics Evolution

The validation loss graph (Figure 4.10) and corresponding training loss graph (Figure 4.11) provide a visual summary of the network's optimization process across 300 epochs. Both the bounding box loss (train/box_loss) and objectness loss show a continuous, steady decay from initial values of 0.091 and 0.037 to final values of approximately 0.029 and 0.020, respectively. The classification loss (train/cls_loss) drops even more significantly, reaching near-zero values, which confirms the model's high confidence in distinguishing between "Safe" and "Unsafe" cable classes. Most importantly, the validation loss curves mirror these training trends without showing the divergence that typically characterizes overfitting. This synchronization indicates that the customized YOLOv5 architecture successfully generalized the underlying patterns of unstructured urban environments rather than merely memorizing training samples, ensuring reliable performance on unseen road data.

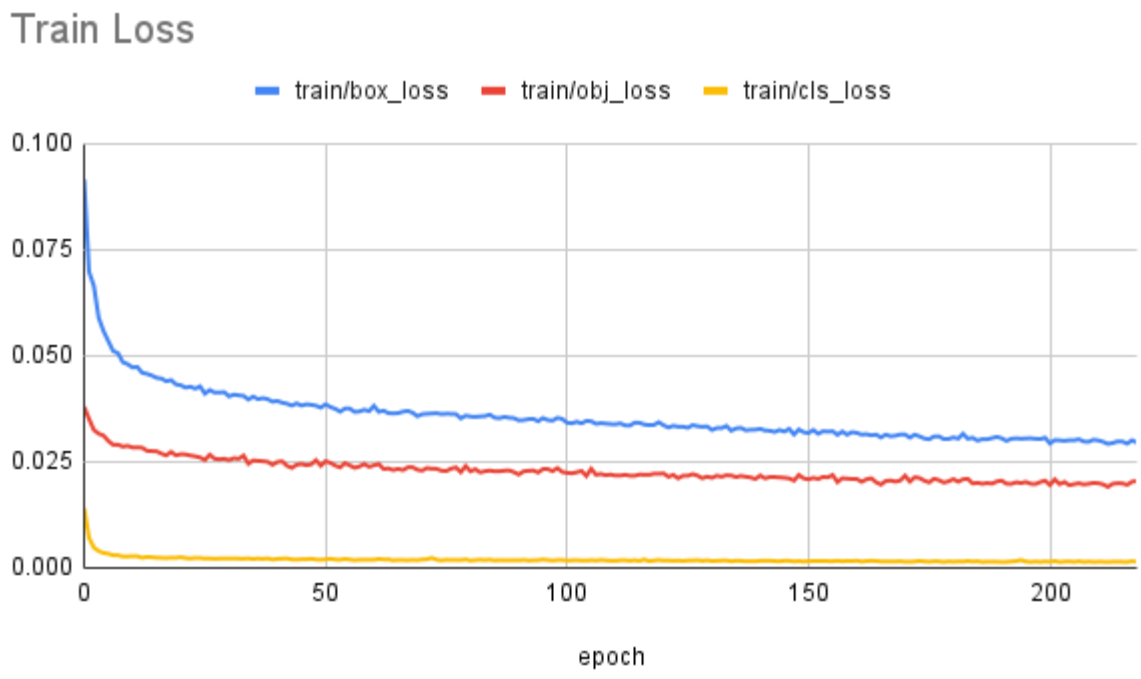

Figure 4.8: Training Loss Convergence

The UFCDD's performance (76% mAP@50-95) sits between these two extremes, demonstrating that the cable's curvilinear, but spatially continuous; nature is less challenging than an unpredictable human pose but far more challenging than a rigid vehicle. This comparative analysis confirms that the localization deficit is a **systemic architectural challenge** for thin and irregular objects, validating the proposed future research direction of integrating specialized loss functions and semantic segmentation principles.

4.4 System Efficiency and Real-Time Feasibility (RQ2 & RQ4)

A critical aspect of the research was ensuring that the detection solution, though specialized, remained viable for safety-critical real-time deployment (RQ2). The customized YOLOv5 model successfully achieved this objective:

- **Low Computational Footprint:** The architectural fine-tuning, reducing the total layers to **157** and parameters to approximately **7 million,** resulted in a low computational load of **15.8 GFLOPs**. This efficiency ensures that the model can be deployed on resource-constrained

AV edge computing platforms, maintaining the minimal latency essential for emergency avoidance maneuvers.

- **Scalability for Smart Mobility (RQ4):** The successful development and validation of the UFCDD and its specialized detection model directly contribute to scalable smart mobility solutions in developing regions [9]. By proving that a high-accuracy detection mechanism can be engineered for a critical regional hazard (fallen cables), the research enhances the global trustworthiness and operational domain of autonomous vehicles. This demonstrates that AV perception systems can be modularly adapted to diverse, real-world infrastructure quality, a necessary step toward achieving global SAE Level 5 autonomy [10].

4.5 Evaluation of Research Questions

The experimental results provide definitive answers to the defined Research Questions, confirming the fulfillment of the research objectives as shown in table 4.2.

Table 4.2: Summary of Research questions and their evaluation

Research Question	Evaluation Summary	Result
RQ1: Data Generation and Fidelity (UFCDD creation)	The collection and annotation of the UFCDD **(2,500+ images from Mirpur, AJK, covering diverse environmental conditions)** successfully provided the specialized, high-fidelity dataset required to train against HSLF cable hazards.	**Validated**

RQ2: Model Architecture and Optimization (YOLOv5 suitability)	The customized YOLOv5 framework proved suitable, achieving a low computational load (**15.8 GFLOPs)** and efficient parameter count **(7 million)**, demonstrating that the architecture meets real-time AV requirements.	**Validated**
RQ3: Performance Validation and Localization Precision (mAP scores)	The model achieved high classification accuracy (**mAP@50: 96%**) but a lower localization score (**mAP@50–95: 76%**). This validates high detection capability while quantifying the inherent geometric difficulty of thin-structure localization.	**Partially Validated (Metrics Quantified)**
RQ4: Impact on Safety and Scalability (Global Relevance)	The framework enhances safety by providing reliable detection for an overlooked HSLF hazard and proves the methodology for scalable adaptation of AV perception to regions with varying infrastructural quality.	**Validated**

Chapter No. 05

CONCLUSION

This research successfully developed a robust, specialized deep learning framework to address the critical safety threat posed by fallen electric and communication cables, an overlooked High-Stakes, Low-Frequency (HSLF) hazard in autonomous vehicle (AV) perception.

The methodology relied on two core accomplishments: the creation of the novel **Urban Fallen Cable Detection Dataset (UFCDD)** (RQ1), and the optimization of a **CNN-based YOLOv5 model** (RQ2). The resulting model demonstrated high operational efficiency, achieving a low computational load of **15.8 GFLOPs,** which is suitable for real-time inference on edge computing platforms.

The performance metrics confirm the framework's capability for hazard identification:

- **High Classification Accuracy:** The model achieved **96% mAP@50,** validating its ability to reliably distinguish hazardous cables from complex background clutters in unstructured urban environments (RQ3).
- **Localization Deficit:** The accuracy dropped to **76% at mAP@50–95**. This result quantifies the primary technical challenge, the failure of conventional bounding box detection to precisely localize thin, irregular, curvilinear objects, confirming a systemic architectural limitation for fine-grained feature resolution.

Ultimately, this study confirms the necessity of developing specialized, domain-specific datasets like UFCDD to ensure scalable and safe autonomous mobility in regions characterized by diverse infrastructural quality (RQ4).

5.1 Future Work

Future efforts must focus on systematically resolving the 76% localization deficit and enhancing the system's robustness for global deployment.

5.1.1 Resolving Geometric and Training Limitations

1. **Architectural Shift to Segmentation:** Transition from bounding box regression to methodologies optimized for pixel-level boundary precision. This involves integrating the UFCDD with frameworks based on **Semantic Segmentation** or **Key-point Encoding** to accurately delineate the cable contour for precise path planning.
2. **Specialized Loss Functions:** Counteract the extreme class imbalance inherent in thin object detection by deploying specialized training methods. This includes augmenting the current loss function with **Focal Loss** and implementing **Boundary-Aware Loss Functions** to

explicitly penalize boundary misalignment errors, thereby forcing the network to refine contour prediction.

5.1.2 Enhancing Robustness and Scalability

1. **Synthetic Data Generation (SDG):** Address data scarcity for HSLF events by leveraging **SDG workflows to** programmatically create high-fidelity, labeled training instances that simulate rare, geometrically complex cable-fall scenarios under diverse conditions.

Joint Detection and Depth Validation: Integrate non-visual confirmation to confirm the physical presence of the hazard. Future work must incorporate **Joint Detection and Depth Estimation to** distinguish hazardous cables lying on the road surface from non-hazardous linear markings or shadows, providing essential 3D safety confirmation.

References:

1. Patil, Ashok Kumar, Bhargav Punugupati, Himanshi Gupta, Niranjan S. Mayur, Srivatsa Ramesh, and Prasad B. Honnavalli. "Building the Future of Transportation: A Comprehensive Survey on AV Perception, Localization, and Mapping." Sensors 25, no. 7 (2025): 2004.
2. Gan, Lu, Wenbo Chu, Guofa Li, Xiaolin Tang, and Keqiang Li. "Large models for intelligent transportation systems and autonomous vehicles: A survey." Advanced Engineering Informatics 62 (2024): 102786.
3. Gherardini, Luca, and Giacomo Cabri. "The impact of autonomous vehicles on safety, economy, society, and environment." World Electric Vehicle Journal 15, no. 12 (2024): 579.
4. Valdovinos, Samantha, Claire Barbrack, and Nicole Podest. "Are Autonomous Vehicles Really "Autonomous?": The Importance of Human Factors in AV Safety." In Proceedings of the Human Factors and Ergonomics Society Annual Meeting, vol. 67, no. 1, pp. 1602-1605. Sage CA: Los Angeles, CA: SAGE Publications, 2023.
5. Valverde, Miguel, Alexandra Moutinho, and João-Vitor Zacchi. "A Survey of Deep Learning-Based 3D Object Detection Methods for Autonomous Driving Across Different Sensor Modalities." Sensors (Basel, Switzerland) 25, no. 17 (2025): 5264.
6. Yan, Huan, and Yong Li. "Generative AI for Intelligent Transportation Systems: Road Transportation Perspective." ACM Computing Surveys (2025).
7. Contreras, Marcelo, Aayush Jain, Neel P. Bhatt, Arunava Banerjee, and Ehsan Hashemi. "A survey on 3D object detection in real time for autonomous driving." Frontiers in Robotics and AI 11 (2024): 1212070.
8. Chen, Zhige, Zhigang Zhang, Qizheng Su, Kai Yang, Yandong Wu, Lei He, and Xiaolin Tang. "Object detection for autonomous vehicles under adverse weather conditions." Expert Systems with Applications (2025): 128994.
9. Yan, Guangjian, Chaoyang Li, Guoqing Zhou, Wuming Zhang, and Xiaowen Li. "Automatic extraction of power lines from aerial images." IEEE Geoscience and remote sensing letters 4, no. 3 (2007): 387-391.
10. Liu, Chenxi. "Situation-aware Customized Machine Intelligence for Transportation Safety, Equity, and Resilience." PhD diss., University of Washington, 2024.
11. Hassija, Vikas, Balamurugan Palanisamy, Arpita Chatterjee, Arpita Mandal, Debanshi Chakraborty, Amit Pandey, G. S. S. Chalapathi, and Dhruv Kumar. "Transformers for Vision: A Survey on Innovative Methods for Computer Vision." IEEE Access (2025).

12. Bozkir, Efe, Süleyman Özdel, Mengdi Wang, Brendan David-John, Hong Gao, Kevin Butler, Eakta Jain, and Enkelejda Kasneci. "Eye-tracked virtual reality: A comprehensive survey on methods and privacy challenges." arXiv preprint arXiv:2305.14080 (2023).
13. Garcia, F., Cerri, P., & Broggi, A. (2017). "Sensor Fusion for Autonomous Vehicles in Urban Environments." IEEE Robotics and Automation Letters, 3(4), 3214-3221. DOI: 10.1109/LRA.2018.2849491.
14. Johnson, R.E., & Thompson, M.K. (2022). "Enhancing Autonomous Driving Algorithms through Visualization of Complex Datasets." Advanced Computing Research, 7(2), 344-360.
15. Gupta, S., & Kumar, P. (2021). Enhancing Low-Light Videos for Autonomous Driving Systems: A Deep Learning Approach. Pattern Recognition Letters.
16. Cordts, M., Omran, M., Ramos, S., Rehfeld, T., Enzweiler, M., Benenson, R., Franke, U., Roth, S., & Schiele, B. (2016). "The Cityscapes Dataset for Semantic Urban Scene Understanding." Proc. of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 3213-3223.
17. Patel, R., & Sengupta, R. (2020). "Autonomous Vehicles: Navigating the Regulatory and Infrastructure Challenges in Developing Nations." Journal of Urban Planning and Development, 146(1), 05019019.
18. Faisal, Md Ahasan Atick, Imene Mecheter, Yazan Qiblawey, Javier Hernandez Fernandez, Muhammad EH Chowdhury, and Serkan Kiranyaz. "Deep learning in automated power line inspection: A review." Applied Energy 385 (2025): 125507.
19. Omrani, Farzane, Hossein Etemadfard, and Rouzbeh Shad. "Assessment of aviation accident datasets in severity prediction through machine learning." Journal of Air Transport Management 115 (2024): 102531.
20. Mohammed, Anaam, Miaad K. Alkhudhairy, and Ali A. Aldhalemi. "Fatal, Non-Fatal Injuries, and Deaths Resulting from Traffic Accidents for the Years 2018-2023 in Iraq." Medical Science Journal for Advance Research 6, no. 1 (2025).
21. Karami, Omid. "3D object detection via LiDAR-camera fusion in autonomous driving." (2024).
22. Dalal, Navneet, and Bill Triggs. "Histograms of oriented gradients for human detection." In 2005 IEEE computer society conference on computer vision and pattern recognition (CVPR'05), vol. 1, pp. 886-893. Ieee, 2005.
23. Ross, T-YLPG, and G. K. H. P. Dollár. "Focal loss for dense object detection." In proceedings of the IEEE conference on computer vision and pattern recognition, pp. 2980-2988. IEEE, 2017.
24. Janai, Joel, Fatma Güney, Aseem Behl, and Andreas Geiger. "Computer vision for autonomous vehicles: Problems, datasets and state of the art." Foundations and trends® in computer graphics and vision 12, no. 1–3 (2020): 1-308.

25. Grigorescu, Sorin, Bogdan Trasnea, Tiberiu Cocias, and Gigel Macesanu. "A survey of deep learning techniques for autonomous driving." Journal of field robotics 37, no. 3 (2020): 362-386.
26. Zhu, Hufei, Yonghui Huang, Ying Xu, Jianhong Zhou, Fuqin Deng, and Yikui Zhai. "Unmanned aerial vehicle (UAV) object detection algorithm based on keypoints representation and rotated distance-IoU loss." Journal of Real-Time Image Processing 21, no. 2 (2024): 58.
27. Liu, Songwen, Hao He, Zhichao Zhang, and Yatong Zhou. "LI-YOLO: An Object Detection Algorithm for UAV Aerial Images in Low-Illumination Scenes." Drones 8, no. 11 (2024): 653.
28. Zhao, Zhong-Qiu, Peng Zheng, Shou-tao Xu, and Xindong Wu. "Object detection with deep learning: A review." IEEE transactions on neural networks and learning systems 30, no. 11 (2019): 3212-3232.
29. Liu, Wei, Dragomir Anguelov, Dumitru Erhan, Christian Szegedy, Scott Reed, Cheng-Yang Fu, and Alexander C. Berg. "Ssd: Single shot multibox detector." In European conference on computer vision, pp. 21-37. Cham: Springer International Publishing, 2016.
30. Sandler, Mark, Andrew Howard, Menglong Zhu, Andrey Zhmoginov, and Liang-Chieh Chen. "Mobilenetv2: Inverted residuals and linear bottlenecks." In Proceedings of the IEEE conference on computer vision and pattern recognition, pp. 4510-4520. 2018.
31. Tremblay, Jonathan, Aayush Prakash, David Acuna, Mark Brophy, Varun Jampani, Cem Anil, Thang To, Eric Cameracci, Shaad Boochoon, and Stan Birchfield. "Training deep networks with synthetic data: Bridging the reality gap by domain randomization." In Proceedings of the IEEE conference on computer vision and pattern recognition workshops, pp. 969-977. 2018.
32. Wigness, Maggie, Sungmin Eum, John G. Rogers, David Han, and Heesung Kwon. "A rugd dataset for autonomous navigation and visual perception in unstructured outdoor environments." In 2019 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS), pp. 5000-5007. IEEE, 2019.
33. Ison, David C. "Comparative analysis of accident and non-accident pilots." Journal of Aviation Technology and Engineering 4, no. 2 (2015): 20.
34. Huang, Chenyu. "Analysis of part 135 aircraft accidents to facilitate flight data monitoring." International Journal of Aviation, Aeronautics, and Aerospace 7, no. 1 (2020).
35. Caesar, Holger, Varun Bankiti, Alex H. Lang, Sourabh Vora, Venice Erin Liong, Qiang Xu, Anush Krishnan, Yu Pan, Giancarlo Baldan, and Oscar Beijbom. "nuscenes: A multimodal dataset for autonomous driving." In Proceedings of the IEEE/CVF conference on computer vision and pattern recognition, pp. 11621-11631. 2020.

36. Azad, Reza, Moein Heidary, Kadir Yilmaz, Michael Hüttemann, Sanaz Karimijafarbigloo, Yuli Wu, Anke Schmeink, and Dorit Merhof. "Loss functions in the era of semantic segmentation: A survey and outlook." arXiv preprint arXiv:2312.05391 (2023).
37. Bouguettaya, Abdelmalek, Ahmed Kechıda, and Amine Mohammed Taberkıt. "A survey on lightweight CNN-based object detection algorithms for platforms with limited computational resources." International Journal of Informatics and Applied Mathematics 2, no. 2 (2019): 28-44.
38. Choe, Tae Eun, Jane Wu, Xiaolin Lin, Karen Kwon, and Minwoo Park. "Hazardnet: Road debris detection by augmentation of synthetic models." In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, pp. 161-171. 2023.
39. Chong, Qianpeng, Mengying Ni, Jianjun Huang, Guangyi Wei, Ziyi Li, and Jindong Xu. "Let the loss impartial: a hierarchical unbiased loss for small object segmentation in high-resolution remote sensing images." European Journal of Remote Sensing 56, no. 1 (2023): 2254473.
40. Damodaran, Satheeswari, Leninisha Shanmugam, and N. M. Jothi Swaroopan. "Overhead power line detection from aerial images using segmentation approaches." Automatika: časopis za automatiku, mjerenje, elektroniku, računarstvo i komunikacije 65, no. 1 (2024): 261-288.
41. Wigness, Maggie, Sungmin Eum, John G. Rogers, David Han, and Heesung Kwon. "A rugd dataset for autonomous navigation and visual perception in unstructured outdoor environments." In 2019 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS), pp. 5000-5007. IEEE, 2019.
42. Qiao, Wenting, Biao Ma, Qiangwei Liu, Xiaoguang Wu, and Gang Li. "Computer vision-based bridge damage detection using deep convolutional networks with expectation maximum attention module." Sensors 21, no. 3 (2021): 824.
43. Gao, Xiaofei. Deep Learning-Based Road Spill Detection Technology. No. 2025-01-7130. SAE Technical Paper, 2025.
44. Gebresenbet, Girma. "Prospects for Implementation of Autonomous Vehicles and Associated Infrastructure in Developing Countries." (2024).
45. Kolbeinsson, Benedikt, and Krystian Mikolajczyk. "UCorr: Wire Detection and Depth Estimation for Autonomous Drones." In International Conference on Robotics, Computer Vision and Intelligent Systems, pp. 179-192. Cham: Springer Nature Switzerland, 2024.
46. Ren, Yuemei, and Lei Li. "Line segment detection algorithm in image extraction improvement study." Journal of Measurements in Engineering 12, no. 1 (2024): 199-213.
47. Boroujeni, Sayed Pedram Haeri, Niloufar Mehrabi, Hazim Alzorgan, Ahmad Sarlak, Mahlagha Fazeli, and Abolfazl Razi. "All You Need for Object Detection: From Pixels, Points, and Prompts

to Next-Gen Fusion and Multimodal LLMs/VLMs in Autonomous Vehicles." arXiv preprint arXiv:2510.26641 (2025).

48. Liew, Jun Hao, Scott Cohen, Brian Price, Long Mai, and Jiashi Feng. "Deep interactive thin object selection." In Proceedings of the IEEE/CVF winter conference on applications of computer vision, pp. 305-314. 2021.
49. Huang, Li, Cheng Chen, Juntong Yun, Ying Sun, Jinrong Tian, Zhiqiang Hao, Hui Yu, and Hongjie Ma. "Multi-scale feature fusion convolutional neural network for indoor small target detection." Frontiers in Neurorobotics 16 (2022): 881021.
50. Madaan, Ratnesh, Daniel Maturana, and Sebastian Scherer. "Wire detection using synthetic data and dilated convolutional networks for unmanned aerial vehicles." In 2017 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS), pp. 3487-3494. IEEE, 2017.
51. Moltz, Jan Hendrik, Annika Hänsch, Bianca Lassen-Schmidt, Benjamin Haas, A. Genghi, J. Schreier, Tomasz Morgas, and Jan Klein. "Learning a loss function for segmentation: A feasibility study." In 2020 IEEE 17th International Symposium on Biomedical Imaging (ISBI), pp. 357-360. IEEE, 2020.
52. O'Mahony, Niall, Sean Campbell, Anderson Carvalho, Suman Harapanahalli, Gustavo Velasco Hernandez, Lenka Krpalkova, Daniel Riordan, and Joseph Walsh. "Deep learning vs. traditional computer vision." In Science and information conference, pp. 128-144. Cham: Springer International Publishing, 2019.
53. Ogundairo, Obaloluwa, Marty Gregg, Kristina Holsgrove, Ying Li, and Mark Allan Kaplan. "A Deep Dive into Boundary-Aware Loss Functions for Fine-Grained Semantic Segmentation." (2024).
54. Woo, Sanghyun, Shoubhik Debnath, Ronghang Hu, Xinlei Chen, Zhuang Liu, In So Kweon, and Saining Xie. "Convnext v2: Co-designing and scaling convnets with masked autoencoders." In Proceedings of the IEEE/CVF conference on computer vision and pattern recognition, pp. 16133-16142. 2023.
55. Sang, Shengtian, Yuyin Zhou, Md Tauhidul Islam, and Lei Xing. "Small-object sensitive segmentation using across feature map attention." IEEE transactions on pattern analysis and machine intelligence 45, no. 5 (2022): 6289-6306.
56. Chaturvedee, Adarsh, Raghad Al-Shabandar, and Ammar H. Mohammed. "Small Object Detection in Autonomous Cars Using a Deep Learning." International Journal of Data Science and Advanced Analytics 6, no. 6 (2024): 307-314.
57. Sun, X., Y. Hou, and P. W. T. Pong. "Underground Power Cable Detection and Inspection Technology Based on Magnetic Field Sensing and Source Reconstruction." In International Symposium on Advanced Magnetic Materials and Applications, ISAMMA 2013. 2013.

58. Wang, Kaige, Tianming Wang, Jianchuang Qu, Huatao Jiang, Qing Li, and Lin Chang. "An end-to-end cascaded image deraining and object detection neural network." IEEE Robotics and Automation Letters 7, no. 4 (2022): 9541-9548.
59. Wang, Xiaoyan, and Huizong Li. "Investigation of Deep Learning Based Semantic Segmentation Models for Autonomous Vehicles." International Journal of Advanced Computer Science & Applications 14, no. 11 (2023).
60. Yi, Wenxin, Zhang Zhang, and Liang Chang. "M4MLF-YOLO: A Lightweight Semantic Segmentation Framework for Spacecraft Component Recognition." Remote Sensing 17, no. 18 (2025): 3144.
61. Elharrouss, Omar, Yasir Mahmood, Yassine Bechqito, Mohamed Adel Serhani, Elarbi Badidi, Jamal Riffi, and Hamid Tairi. "Loss functions in deep learning: a comprehensive review." arXiv e-prints (2025): arXiv-2504.
62. Liu, Jun, Zhenglun Kong, Pu Zhao, Weihao Zeng, Hao Tang, Xuan Shen, Changdi Yang et al. "Tsla: A task-specific learning adaptation for semantic segmentation on autonomous vehicles platform." IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems (2024).
63. Zhang, Yijie, Zhong Liu, Hui Nie, Yuechen Yu, and Guannan Chang. "Vehicle target detection based on ConvNext-improved YoloV5." In International Conference on Image Processing and Deep Learning (IPDL 2025), vol. 13707, pp. 163-169. SPIE, 2025.
64. Zhou, Chen, Jiaolong Yang, Chunshui Zhao, and Gang Hua. "Fast, accurate thin-structure obstacle detection for autonomous mobile robots." In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshops, pp. 1-10. 2017.
65. Gwizdała, Jakub, Doruk Oner, Soumava Kumar Roy, Mian Akbar Shah, Ad Eberhard, Ivan Egorov, Philipp Krüsi, Grigory Yakushev, and Pascal Fua. "Vision-based power line cables and pylons detection for low flying aircraft." Machine Vision and Applications 36, no. 2 (2025): 46.